

## Item 1

### git\_comments:

1. \* \* @return The length in bytes of the on disk size for this SSTable. For \* compressed files, this is not the same thing as the data length (see \* length())
2. This differs from length for compressed files (but we still need length for SegmentIterator because offsets in the file are relative to the uncompressed size)

### git\_commits:

1. **summary:** Fix places where uncompressed sstable size is used in place of the compressed one.  
**message:** Fix places where uncompressed sstable size is used in place of the compressed one. patch by slebresne; reviewed by jbellis for CASSANDRA-3338 git-svn-id: <https://svn.apache.org/repos/asf/cassandra/branches/cassandra-1.0.0@1180970> 13f79535-47bb-0310-9956-ffa450edef68

### github\_issues:

### github\_issues\_comments:

### github\_pulls:

### github\_pulls\_comments:

### github\_pulls\_reviews:

### jira\_issues:

1. **summary:** Uncompressed sizes are used to estimate space for compaction of compressed sstables  
**description:** We are using the uncompressed data size when estimating if we have enough to compact sstables. This means we can easily refuse compaction when there is clearly enough room to compact.

### jira\_issues\_comments:

1. maybe we should rename .length() to uncompressedLength() as well? otherwise +1
2. **body:** Actually renaming length() to uncompressedLength() proved to be a good idea, as I had missed quite a bunch of places where the onDiskLength should be used. Attached v2.  
**label:** code-design
3. +1
4. Committed, thanks