

Item 8

**git\_comments:**

**git\_commits:**

1. **summary:** HADOOP-2339 Delete command with no WHERE clause  
**message:** HADOOP-2339 Delete command with no WHERE clause git-svn-id:  
<https://svn.apache.org/repos/asf/lucene/hadoop/trunk/src/contrib/hbase@600771> 13f79535-47bb-0310-9956-ffa450edef68

**github\_issues:**

**github\_issues\_comments:**

**github\_pulls:**

**github\_pulls\_comments:**

**github\_pulls\_reviews:**

**jira\_issues:**

1. **summary:** Delete command with no WHERE clause  
**description:** using HbaseAdmin.deleteColumn() method. {code} DELETE column\_name FROM table\_name; {code}

**jira\_issues\_comments:**

1. {code} hql > select \* from udanax; +-----+-----+-----+ |  
Row | Column | Cell | +-----+-----+-----+ | row1 | a: | aa |  
+-----+-----+-----+ | row1 | b: | aa | +-----+-----+  
-+-----+-----+ | row2 | b: | aa | +-----+-----+  
--+-----+ 3 row(s) in set. (0.06 sec) hql > delete b: from udanax; 07/12/04 09:30:15 INFO  
hbase.HBaseAdmin: Disabled table udanax 07/12/04 09:30:15 INFO hbase.HBaseAdmin: Enabled table  
udanax Column(s) deleted successfully. (10.16 sec) hql > select \* from udanax; +-----+-----+  
-----+-----+ | Row | Column | Cell | +-----+-----+  
-----+-----+ | row1 | a: | aa | +-----+-----+  
-----+ 1 row(s) in set. (0.05 sec) hql > {code}
2. submitting. I'd like to commit to 0.16 TRUNK.
3. {code} DELETE \* FROM table\_name; {code} NOTE: Now just using deleteColumn() method in a loop.
4. **body:** +1 overall. Here are the results of testing the latest attachment  
<http://issues.apache.org/jira/secure/attachment/12370898/2339.patch> against trunk revision r600707.  
@author +1. The patch does not contain any @author tags. javadoc +1. The javadoc tool did not generate  
any warning messages. javac +1. The applied patch does not generate any new compiler warnings.  
findbugs +1. The patch does not introduce any new Findbugs warnings. core tests +1. The patch passed  
core unit tests. contrib tests +1. The patch passed contrib unit tests. Test results:  
<http://lucene.zones.apache.org:8080/hudson/job/Hadoop-Patch/1250/testReport/> Findbugs warnings:  
<http://lucene.zones.apache.org:8080/hudson/job/Hadoop-Patch/1250/artifact/trunk/build/test/findbugs/newPatchFindbugsWarnings.html> Checkstyle results:  
<http://lucene.zones.apache.org:8080/hudson/job/Hadoop-Patch/1250/artifact/trunk/build/test/checkstyle-errors.html> Console output: <http://lucene.zones.apache.org:8080/hudson/job/Hadoop-Patch/1250/console>  
This message is automatically generated.  
**label:** code-design
5. Committed. Resolving. Thanks Edward.
6. One question: What happens if user specifies a column name that is not a column family? I suppose in  
this case we should delete all cells for the specified column? Also, I wonder if 'delete b: from udanax'  
shouldn't just delete column content rather than actual column itself: i.e. disable, remove, readd, re-  
enable?

7. Ah... Sorry for my carelessness. I'll attach the patch for describe below. {code} DELETE column\_name:qualifier\_name FROM table\_name; {code}
8. ugh... It seems troublesome... \*hql > DELETE columnA: FROM table;\* In this case, only "columnA:" should be deleted. {code} Row1 columnA: ----> to be deleted columnA:a columnA:b Row2 columnA: ---> to be deleted columnA:b columnA:c {code} .... And then, how can i delete all columns in columnfamily? I'd like to use the deleteColumn() method to drop the \*ColumnFamily\* {code} DROP TABLE table\_name; DROP COLUMNFAMILY columnfamily\_name; {code}
9. {code} DROP COLUMNFAMILY columnfamily\_name1[, columnfamily\_name2, ...] FROM table\_name; {code}
10. Ok. Drop has a different mean. So,... column\_name1[, column\_name2, ...] or COLUMNFAMILIES(column1[, column2]) ex. DELETE column:, column:a, column:b FROM table; // delete the specified columns DELETE COLUMNFAMILIES(column1, column2) FROM table; // delete the columnfamilies DELETE \* FROM table; {code} hql > select \* from udanax; +-----+  
+-----+ | Row | Column | Cell | +-----+  
+-----+ | a | b: | bb3 | +-----+  
+-----+ | a | c: | cc | +-----+ | b | b: | bb2 | +-----+  
+-----+ | b | c: | cc2 | +-----+  
+-----+ | c | b:aa | aa | +-----+  
+-----+ | c | c: | bb | +-----+ 6 row(s)  
in set. (0.06 sec) hql > delete 'b:aa' from udanax; 1 Column(s) deleted successfully. (0.03 sec) hql > select \* from udanax; +-----+  
+-----+ | Row | Column | Cell | +-----+  
+-----+ | a | b: | bb3 | +-----+  
+-----+ | a | c: | cc | +-----+  
+-----+ | b | b: | bb2 | +-----+ | b |  
+-----+ | c: | cc2 | +-----+ | c | c: | bb | +-----+  
+-----+ 5 row(s) in set. (0.06 sec) hql > delete  
columnfamilies(c) from udanax; delete columnfamilies(c) from udanax; Dropping c:... Please wait.  
07/12/04 16:36:38 INFO hbase.HBaseAdmin: Disabled table udanax 07/12/04 16:36:38 INFO  
hbase.HBaseAdmin: Enabled table udanax 1 columnfamily(s) deleted successfully. (10.12 sec) hql >  
select \* from udanax; select \* from udanax; +-----+  
+-----+ | Row | Column | Cell | +-----+ | a | b: |  
bb3 | +-----+ | b | b: | bb2 | +-----+  
+-----+ 2 row(s) in set. (0.07 sec) hql > {code}
11. **body:** Sorry, wrong spell was fixed. (dropped -> deleted)  
**label:** documentation
12. Ok, submitting.
13. -1 overall. Here are the results of testing the latest attachment  
[http://issues.apache.org/jira/secure/attachment/12370921/2339\\_v03.patch](http://issues.apache.org/jira/secure/attachment/12370921/2339_v03.patch) against trunk revision r600771.  
@author +1. The patch does not contain any @author tags. javadoc +1. The javadoc tool did not generate any warning messages. javac +1. The applied patch does not generate any new compiler warnings.  
findbugs +1. The patch does not introduce any new Findbugs warnings. core tests +1. The patch passed core unit tests. contrib tests -1. The patch failed contrib unit tests. Test results:  
<http://lucene.zones.apache.org:8080/hudson/job/Hadoop-Patch/1255/testReport/> Findbugs warnings:  
<http://lucene.zones.apache.org:8080/hudson/job/Hadoop-Patch/1255/artifact/trunk/build/test/findbugs/newPatchFindbugsWarnings.html> Checkstyle results:  
<http://lucene.zones.apache.org:8080/hudson/job/Hadoop-Patch/1255/artifact/trunk/build/test/checkstyle-errors.html> Console output: <http://lucene.zones.apache.org:8080/hudson/job/Hadoop-Patch/1255/console>  
This message is automatically generated.
14. Local test passed. So, re-submitting. {code} test: [echo] contrib: hbase [delete] Deleting directory /root/workspace/hadoop/build/contrib/hbase/test/logs [mkdir] Created dir: /root/workspace/hadoop/build/contrib/hbase/test/logs [junit] Running org.apache.hadoop.hbase.shell.TestHBaseShell [junit] Tests run: 2, Failures: 0, Errors: 0, Time elapsed: 52.787 sec [junit] Running org.apache.hadoop.hbase.shell.TestSubstitutionVariables [junit] Tests run: 2, Failures: 0, Errors: 0, Time elapsed: 0.019 sec [junit] Running org.apache.hadoop.hbase.shell.algebra.TestBooleanCondition [junit] Tests run: 1, Failures: 0, Errors: 0, Time elapsed: 0.044 sec [junit] Running org.apache.hadoop.hbase.shell.algebra.TestBooleanTermFilter [junit] Tests run: 1, Failures: 0, Errors: 0, Time elapsed: 23.241 sec [junit] Running org.apache.hadoop.hbase.shell.algebra.TestJoinCondition [junit] Tests run: 1, Failures: 0, Errors: 0, Time elapsed: 0.158 sec [junit] Running org.apache.hadoop.hbase.shell.algebra.TestTableJoinMapReduce

- [junit] Tests run: 1, Failures: 0, Errors: 0, Time elapsed: 87.233 sec [junit] Running org.apache.hadoop.hbase.util.TestBase64 [junit] Tests run: 1, Failures: 0, Errors: 0, Time elapsed: 0.136 sec [junit] Running org.apache.hadoop.hbase.util.TestKeying [junit] Tests run: 1, Failures: 0, Errors: 0, Time elapsed: 0.011 sec [junit] Running org.onelab.test.TestFilter [junit] Tests run: 3, Failures: 0, Errors: 0, Time elapsed: 0.025 sec BUILD SUCCESSFUL Total time: 2 minutes 54 seconds bash-3.00# {code}
15. Re-try
16. -1 overall. Here are the results of testing the latest attachment [http://issues.apache.org/jira/secure/attachment/12370926/2339\\_v04.patch](http://issues.apache.org/jira/secure/attachment/12370926/2339_v04.patch) against trunk revision r600771. @author +1. The patch does not contain any @author tags. javadoc +1. The javadoc tool did not generate any warning messages. javac +1. The applied patch does not generate any new compiler warnings. findbugs +1. The patch does not introduce any new Findbugs warnings. core tests +1. The patch passed core unit tests. contrib tests -1. The patch failed contrib unit tests. Test results: <http://lucene.zones.apache.org:8080/hudson/job/Hadoop-Patch/1256/testReport/> Findbugs warnings: <http://lucene.zones.apache.org:8080/hudson/job/Hadoop-Patch/1256/artifact/trunk/build/test/findbugs/newPatchFindbugsWarnings.html> Checkstyle results: <http://lucene.zones.apache.org:8080/hudson/job/Hadoop-Patch/1256/artifact/trunk/build/test/checkstyle-errors.html> Console output: <http://lucene.zones.apache.org:8080/hudson/job/Hadoop-Patch/1256/console> This message is automatically generated.
17. Integrated in Hadoop-Nightly #322 (See [<http://lucene.zones.apache.org:8080/hudson/job/Hadoop-Nightly/322/>])
18. **body:** -1 overall. Here are the results of testing the latest attachment [http://issues.apache.org/jira/secure/attachment/12370926/2339\\_v04.patch](http://issues.apache.org/jira/secure/attachment/12370926/2339_v04.patch) against trunk revision r600771. @author +1. The patch does not contain any @author tags. javadoc +1. The javadoc tool did not generate any warning messages. javac +1. The applied patch does not generate any new compiler warnings. findbugs +1. The patch does not introduce any new Findbugs warnings. core tests +1. The patch passed core unit tests. contrib tests -1. The patch failed contrib unit tests. Test results: <http://lucene.zones.apache.org:8080/hudson/job/Hadoop-Patch/1257/testReport/> Findbugs warnings: <http://lucene.zones.apache.org:8080/hudson/job/Hadoop-Patch/1257/artifact/trunk/build/test/findbugs/newPatchFindbugsWarnings.html> Checkstyle results: <http://lucene.zones.apache.org:8080/hudson/job/Hadoop-Patch/1257/artifact/trunk/build/test/checkstyle-errors.html> Console output: <http://lucene.zones.apache.org:8080/hudson/job/Hadoop-Patch/1257/console> This message is automatically generated.  
**label:** code-design
19. Canceling because QA returns -1
20. Re-trying after 2333 comitted.
21. -1 overall. Here are the results of testing the latest attachment [http://issues.apache.org/jira/secure/attachment/12370926/2339\\_v04.patch](http://issues.apache.org/jira/secure/attachment/12370926/2339_v04.patch) against trunk revision r601232. @author +1. The patch does not contain any @author tags. javadoc +1. The javadoc tool did not generate any warning messages. javac +1. The applied patch does not generate any new compiler warnings. findbugs +1. The patch does not introduce any new Findbugs warnings. core tests +1. The patch passed core unit tests. contrib tests -1. The patch failed contrib unit tests. Test results: <http://lucene.zones.apache.org:8080/hudson/job/Hadoop-Patch/1270/testReport/> Findbugs warnings: <http://lucene.zones.apache.org:8080/hudson/job/Hadoop-Patch/1270/artifact/trunk/build/test/findbugs/newPatchFindbugsWarnings.html> Checkstyle results: <http://lucene.zones.apache.org:8080/hudson/job/Hadoop-Patch/1270/artifact/trunk/build/test/checkstyle-errors.html> Console output: <http://lucene.zones.apache.org:8080/hudson/job/Hadoop-Patch/1270/console> This message is automatically generated.
22. To clarify this issue, are we talking about what is essentially an ALTER TABLE DROP COLUMN in SQL? If so, the description should be changed to reflect that.
23. Not a major issue, but should still be looked at.
24. canceling
25. resubmitting
26. +1 overall. Here are the results of testing the latest attachment [http://issues.apache.org/jira/secure/attachment/12370926/2339\\_v04.patch](http://issues.apache.org/jira/secure/attachment/12370926/2339_v04.patch) against trunk revision r603084. @author +1. The patch does not contain any @author tags. javadoc +1. The javadoc tool did not generate any warning messages. javac +1. The applied patch does not generate any new compiler warnings. findbugs +1. The patch does not introduce any new Findbugs warnings. core tests +1. The patch passed core unit tests. contrib tests +1. The patch passed contrib unit tests. Test results: <http://lucene.zones.apache.org:8080/hudson/job/Hadoop-Patch/1314/testReport/> Findbugs warnings:

<http://lucene.zones.apache.org:8080/hudson/job/Hadoop-Patch/1314/artifact/trunk/build/test/findbugs/newPatchFindbugsWarnings.html> Checkstyle results:  
<http://lucene.zones.apache.org:8080/hudson/job/Hadoop-Patch/1314/artifact/trunk/build/test/checkstyle-errors.html> Console output: <http://lucene.zones.apache.org:8080/hudson/job/Hadoop-Patch/1314/console>  
This message is automatically generated.

27. Now DELETE statement has a serious logical error, It should be fixed, and I think it is a important thing.
28. **body:** Let's specify very clearly what this issue regards. As I understand it, the goal is to be able to delete a specific column family and qualifier on every row in the entire table. It's important to note that the qualifier is needed, because you can already do this essentially with ALTER TABLE if you want to delete an entire column family. Is this correct? If so, I'm not sure that we should offer this functionality. It will be slow, especially if it's run as a single delete at a time in serial from the shell. I would think that from the shell, you have the ability to empty an entire table with TRUNCATE, or delete an entire column family with ALTER TABLE ... DROP columnfamily, and that's probably enough. If you actually do need to delete just a given qualified column name on every row in your table, you should write a mapreduce job or some other highly parallel way to do that. This is in keeping with the goal of making the shell both efficient and lean when it comes to the functionality it supports.  
**label:** code-design
29. **body:** 'columnfamily', 'columnfamily:specified\_qualifer' delete function is important thing. I'm not bothering to make a example. Also, I thought it should be work as inter-row wise partitioned sub structre parallelism. But, now shell is 0.0.1 version. Imagine the processing-speed:increment by row-count:increment 2d plots of scanner and mapreduce, start points of processing-speed dimension and slopes will be different. So, I thought the cross point find algorithm for the most faster automatical processing. 'Slow Performance' -can't- should not touch the grammar rules.  
**label:** code-design
30. Let's assume the webtable (bigtable paper) {code} url<row> content anchor language metatag ...  
<http://abcd.com> ... anchor:<http://xyz.com> <xyz dot com> ... ... anchor:<http://udanax.org> <udanax homepage> {code} At this time, 'udanax.org' was evaporation from the world-wide-web. Then we need to delete the 'anchor:<http://udanax.org>' of all rows for transformation of url by anchor (inlink by outlink).
31. Bryan, I removed your comments because there is no messages. If you have a constructive comment, please share it.
32. Anyhow, Should it be closed?
33. Canceling.
34. This is a new issue, it'll be re-made.
35. I have the same issue. Does this fix locate in branch-0.16 in svn?