

**git\_comments:**

1. replace-overwrite introduces no new files

**git\_commits:**

1. **summary:** HIVE-14268 : INSERT-OVERWRITE is not generating an INSERT event during hive replication (Sushanth Sowmyan, reviewed by Alan Gates)  
**message:** HIVE-14268 : INSERT-OVERWRITE is not generating an INSERT event during hive replication (Sushanth Sowmyan, reviewed by Alan Gates)

**github\_issues:**

**github\_issues\_comments:**

**github\_pulls:**

**github\_pulls\_comments:**

**github\_pulls\_reviews:**

**jira\_issues:**

1. **summary:** INSERT-OVERWRITE is not generating an INSERT event during hive replication  
**description:** During Hive replication invoked from falcon, the source cluster did not generate appropriate INSERT events associated with the INSERT OVERWRITE, generating only an ALTER PARTITION event. However, an ALTER PARTITION is a metadata-only event, and thus, only metadata changes were replicated across, modifying the metadata of the destination, while not updating the data.

**jira\_issues\_comments:**

1. Patch attached.
2. Slightly more complex version in .2.patch, with thrift change to allow signalling on the event as to whether this is an overwrite event or not (although we still don't use that info on the metastore side for now)
3. [~gates], could you please review? (We can go with either approach)
4. Given that there's no use for the replace information on the server, for now I say we go with patch 1. If we find some use for propagating that information in the future we can add it to thrift then. +1 for patch 1.
5. Sounds good - reuploading .1.patch as .3.patch so the tests run on that.
6. reuploading .1/.3.patch as .4.patch because the builds.apache.org job borked and is not picking it up again.
7. Here are the results of testing the latest attachment: <https://issues.apache.org/jira/secure/attachment/12819192/HIVE-14268.4.patch> {color:green}SUCCESS:{color} +1 due to 1 test(s) being added or modified. {color:red}ERROR:{color} -1 due to 10 failed/errored test(s), 10341 tests executed \*Failed tests:\* {noformat} TestMsgBusConnection - did not produce a TEST-\*.xml file org.apache.hadoop.hive.cli.TestCliDriver.testCliDriver\_acid\_globallimit  
org.apache.hadoop.hive.cli.TestCliDriver.testCliDriver\_list\_bucket\_dml\_12  
org.apache.hadoop.hive.cli.TestCliDriver.testCliDriver\_stats\_list\_bucket  
org.apache.hadoop.hive.cli.TestCliDriver.testCliDriver\_subquery\_multiinsert  
org.apache.hadoop.hive.cli.TestMiniTezCliDriver.testCliDriver\_acid\_globallimit  
org.apache.hadoop.hive.llap.daemon.impl.TestLlapTokenChecker.testCheckPermissions  
org.apache.hadoop.hive.llap.daemon.impl.TestLlapTokenChecker.testGetToken  
org.apache.hadoop.hive.llap.tezplugins.TestLlapTaskSchedulerService.testDelayedLocalityNodeCommErrorImmediateAllocation  
org.apache.hadoop.hive.metastore.TestMetaStoreMetrics.testConnections {noformat} Test results:  
<https://builds.apache.org/job/PreCommit-HIVE-MASTER-Build/598/testReport> Console output:  
<https://builds.apache.org/job/PreCommit-HIVE-MASTER-Build/598/console> Test logs: <http://ec2-204-236-174-241.us-west-1.compute.amazonaws.com/logs/PreCommit-HIVE-MASTER-Build-598/> Messages: {noformat} Executing  
org.apache.hive.ptest.execution.TestCheckPhase Executing org.apache.hive.ptest.execution.PreparePhase Executing  
org.apache.hive.ptest.execution.ExecutionPhase Executing org.apache.hive.ptest.execution.ReportingPhase Tests exited with:  
TestsFailedException: 10 tests failed {noformat} This message is automatically generated. ATTACHMENT ID: 12819192 -  
PreCommit-HIVE-MASTER-Build
8. None of the test failures are related, going ahead and committing.
9. Note that this issue was introduced only by <https://issues.apache.org/jira/browse/HIVE-12907> , and thus, does not need to be applied to prior versions unless HIVE-12907 is backported to others.