

Item 69

git_comments:

1. normalize the path
2. TODO sanity check col fam

git_commits:

1. **summary:** ACCUMULO-1747 expanded use of default uri and changed how relative paths are handled in code
message: ACCUMULO-1747 expanded use of default uri and changed how relative paths are handled in code

github_issues:

github_issues_comments:

github_pulls:

github_pulls_comments:

github_pulls_reviews:

jira_issues:

1. **summary:** ServerConstants.getBaseDirs() use fs default name, not instance.dfs.uri
description: I've been trying to track down why I'm seeing connection refused to port 9000 in {{accumulo init}} -looks to me like 2013-10-02 19:04:12,600 [util.Initialize] INFO : Hadoop Filesystem is file:/ 2013-10-02 19:04:12,678 [Configuration.deprecation] INFO : fs.default.name is deprecated. Instead, use fs.defaultFS 2013-10-02 19:04:12,678 [util.Initialize] INFO : Accumulo data dirs are [hdfs://localhost:9000/Users/stevil/.hoya/cluster/TestAcc Initialize is picking up instance.dfs.uri value, but the {{ServerConstants.getBaseDirs()}} is picking up the default FS -not even from accumulo-site.xml by the look of things, so fails to work except when instance.dfs.uri==fs.default.name

jira_issues_comments:

1. generated site file, has both the 1.x and the (new, changed) property for the default fs: {code} <?xml version="1.0" encoding="UTF-8" standalone="no"?><configuration> <property>
<name>trace.user</name><value>root</value>
<source>file:/Users/stevil/.hoya/cluster/TestAccM1T0/original/accumulo-site.xml</source></property>
<property><name>instance.secret</name><value>d4483896-76a7-42a2-a0ad-cf93f44d5305</value>
<source>programatically</source></property> <property>
<name>tserver.memory.maps.native.enabled</name><value>>false</value>
<source>file:/Users/stevil/.hoya/cluster/TestAccM1T0/original/accumulo-site.xml</source></property>
<property><name>tserver.memory.maps.max</name><value>80M</value>
<source>file:/Users/stevil/.hoya/cluster/TestAccM1T0/original/accumulo-site.xml</source></property>
<property><name>tserver.sort.buffer.size</name><value>50M</value>
<source>file:/Users/stevil/.hoya/cluster/TestAccM1T0/original/accumulo-site.xml</source></property>
<property><name>logger.dir.walog</name><value>walog</value>
<source>file:/Users/stevil/.hoya/cluster/TestAccM1T0/original/accumulo-site.xml</source></property>
<property><name>instance.dfs.uri</name><value>file:/</value><source>programatically</source>
</property> <property><name>tserver.walog.max.size</name><value>100M</value>
<source>file:/Users/stevil/.hoya/cluster/TestAccM1T0/original/accumulo-site.xml</source></property>
<property><name>monitor.port.client</name><value>50095</value><source>programatically</source>
</property> <property><name>hoya.template.origin</name>
<value>file:/Users/stevil/.hoya/cluster/TestAccM1T0/original/accumulo-site.xml</value>
<source>programatically</source></property> <property>
<name>trace.token.property.password</name><value>secret</value>
<source>file:/Users/stevil/.hoya/cluster/TestAccM1T0/original/accumulo-site.xml</source></property>
<property><name>instance.zookeeper.host</name><value>127.0.0.1:64731</value>

```

<source>programatically</source></property> <property><name>instance.dfs.dir</name>
<value>/Users/stevel/.hoya/cluster/TestAccM1T0/database</value><source>programatically</source>
</property> <property><name>tserver.cache.index.size</name><value>20M</value>
<source>file:/Users/stevel/.hoya/cluster/TestAccM1T0/original/accumulo-site.xml</source></property>
<property><name>master.port.client</name><value>9999</value><source>programatically</source>
</property> <property><name>fs.defaultFS</name><value>file:///</value>
<source>programatically</source></property> <property><name>tserver.cache.data.size</name>
<value>7M</value><source>file:/Users/stevel/.hoya/cluster/TestAccM1T0/original/accumulo-
site.xml</source></property> <property><name>trace.port.client</name><value>0</value>
<source>programatically</source></property> <property><name>general.classpaths</name><value>
$ACCUMULO_HOME/lib/accumulo-server.jar, $ACCUMULO_HOME/lib/accumulo-core.jar,
$ACCUMULO_HOME/lib/accumulo-start.jar, $ACCUMULO_HOME/lib/accumulo-fate.jar,
$ACCUMULO_HOME/lib/accumulo-proxy.jar, $ACCUMULO_HOME/lib/[^].*.jar,
$ZOOKEEPER_HOME/zookeeper[^].*.jar, $HADOOP_CONF_DIR, $HADOOP_PREFIX/[^].*.jar,
$HADOOP_PREFIX/lib/[^].*.jar, $HADOOP_PREFIX/share/hadoop/common/*.jar,
$HADOOP_PREFIX/share/hadoop/common/lib/*.jar, $HADOOP_PREFIX/share/hadoop/hdfs/*.jar,
$HADOOP_PREFIX/share/hadoop/mapreduce/*.jar, $HADOOP_PREFIX/share/hadoop/yarn/*.jar,
</value><source>file:/Users/stevel/.hoya/cluster/TestAccM1T0/original/accumulo-site.xml</source>
</property> <property><name>fs.default.name</name><value>file:///</value> {code} It
does look like I'm not using file:/// as the full URI, but its moot: ServerConstants doesn't use the instance
uri: {code} hadoopConfig.get("fs.default.name") {code} As to why the FS settings aren't being picked up,
this test used a core-site.xml pointing to hdfs://local:9000 -which got picked up instead of the site.xml
value -steve

```

2. I think fs.defaultFS wasn't introduced until hadoop 2, so this needs to be treaded lightly around.
3. you can get away without the fs.defaultFS stuff as hadoop2 fallls back to "fs.default.name" -I'm just pushing both down for consistency. What is needed here is that the base dir for files ought -if I've understood it- be build from {{instance.dfs.uri + instance.dfs.dir}}, but instead it is {{fs.default.name + instance.dfs.dir}}, so you can't run accumulo against a non-default filesystem
4. I just ran into this too. The code is not using instance.dfs.dir and the new prop instance.volumes consistently. I am going to make the code use these props consistently.
5. Commit 95cb49043ddfe31631639ccfd6bea70201e2b5e4 in branch refs/heads/master from [~keith_turner] [<https://git-wip-us.apache.org/repos/asf?p=accumulo.git;h=95cb490>] ACCUMULO-1747 fixed, defined, and documented usage of intance.dfs.uri and instance.volumes
6. Commit cb45e89b0f78d170cb0f5383ab08dea780c9386e in branch refs/heads/master from [~keith_turner] [<https://git-wip-us.apache.org/repos/asf?p=accumulo.git;h=cb45e89>] ACCUMULO-1747 expanded use of default uri and changed how relative paths are handled in code
7. marking as resolved as the changes are in & I haven't seen this for a while