

git_comments:

1. * * Closes this snapshot releasing any associated resources. Once a * snapshot is closed, no new iterators can be created for it.
2. * * An immutable view of the {@link KeyValueStore} at a point-in-time. * The snapshot MUST be closed after use. * * @param <K> key type * @param <V> value type
3. * * Creates a new iterator for this snapshot. The iterator MUST be * closed after its execution by invoking {@link KeyValueIterator#close}. * @return an iterator
4. * * Closes this snapshot releasing any associated resources. Once a * snapshot is closed, no new iterators can be created for it.
5. * * An immutable view of the {@link KeyValueStore} at a point-in-time. * The snapshot MUST be closed after use. * * @param <K> key type * @param <V> value type
6. * * Creates a new iterator for this snapshot. The iterator MUST be * closed after its execution by invoking {@link KeyValueIterator#close}. * @return an iterator
7. Please comment out range query part in order to do an accurate perf test for snapshot

git_commits:

1. **summary:** SAMZA-1705: Switch to use snapshot in iterable impl of RocksDb
message: SAMZA-1705: Switch to use snapshot in iterable impl of RocksDb We should use rocksDb.snapshot() method to keep the snapshot and creates a new iterator with it all the time. The perf shows a little bit more expensive but mostly on par with range iterator query. Author: xinyuiscool <xiliu@linkedin.com> Reviewers: Jagadish V <vjagadish@apache.org> Closes #510 from xinyuiscool/SAMZA-1705 (cherry picked from commit 89beb1fccb01c781a4de905d57a4bd99df25577a) Signed-off-by: xiliu <xiliu@linkedin.com>

github_issues:**github_issues_comments:****github_pulls:**

1. **title:** SAMZA-1705: Switch to use snapshot in iterable impl of RocksDb
body: We should use rocksDb.snapshot() method to keep the snapshot and creates a new iterator with it all the time. The perf shows a little bit more expensive but mostly on par with range iterator query.
label: code-design
2. **title:** SAMZA-1705: Switch to use snapshot in iterable impl of RocksDb
body: We should use rocksDb.snapshot() method to keep the snapshot and creates a new iterator with it all the time. The perf shows a little bit more expensive but mostly on par with range iterator query.
3. **title:** SAMZA-1705: Switch to use snapshot in iterable impl of RocksDb
body: We should use rocksDb.snapshot() method to keep the snapshot and creates a new iterator with it all the time. The perf shows a little bit more expensive but mostly on par with range iterator query.
4. **title:** SAMZA-1705: Switch to use snapshot in iterable impl of RocksDb
body: We should use rocksDb.snapshot() method to keep the snapshot and creates a new iterator with it all the time. The perf shows a little bit more expensive but mostly on par with range iterator query.
5. **title:** SAMZA-1705: Switch to use snapshot in iterable impl of RocksDb
body: We should use rocksDb.snapshot() method to keep the snapshot and creates a new iterator with it all the time. The perf shows a little bit more expensive but mostly on par with range iterator query.
label: documentation
6. **title:** SAMZA-1705: Switch to use snapshot in iterable impl of RocksDb
body: We should use rocksDb.snapshot() method to keep the snapshot and creates a new iterator with it all the time. The perf shows a little bit more expensive but mostly on par with range iterator query.
7. **title:** SAMZA-1705: Switch to use snapshot in iterable impl of RocksDb
body: We should use rocksDb.snapshot() method to keep the snapshot and creates a new iterator with it all the time. The perf shows a little bit more expensive but mostly on par with range iterator query.
label: documentation
8. **title:** SAMZA-1705: Switch to use snapshot in iterable impl of RocksDb

body: We should use `rocksDb.snapshot()` method to keep the snapshot and creates a new iterator with it all the time. The perf shows a little bit more expensive but mostly on par with range iterator query.

label: code-design

9. **title:** SAMZA-1705: Switch to use snapshot in iterable impl of RocksDb
body: We should use `rocksDb.snapshot()` method to keep the snapshot and creates a new iterator with it all the time. The perf shows a little bit more expensive but mostly on par with range iterator query.
10. **title:** SAMZA-1705: Switch to use snapshot in iterable impl of RocksDb
body: We should use `rocksDb.snapshot()` method to keep the snapshot and creates a new iterator with it all the time. The perf shows a little bit more expensive but mostly on par with range iterator query.
11. **title:** SAMZA-1705: Switch to use snapshot in iterable impl of RocksDb
body: We should use `rocksDb.snapshot()` method to keep the snapshot and creates a new iterator with it all the time. The perf shows a little bit more expensive but mostly on par with range iterator query.
12. **title:** SAMZA-1705: Switch to use snapshot in iterable impl of RocksDb
body: We should use `rocksDb.snapshot()` method to keep the snapshot and creates a new iterator with it all the time. The perf shows a little bit more expensive but mostly on par with range iterator query.
13. **title:** SAMZA-1705: Switch to use snapshot in iterable impl of RocksDb
body: We should use `rocksDb.snapshot()` method to keep the snapshot and creates a new iterator with it all the time. The perf shows a little bit more expensive but mostly on par with range iterator query.
14. **title:** SAMZA-1705: Switch to use snapshot in iterable impl of RocksDb
body: We should use `rocksDb.snapshot()` method to keep the snapshot and creates a new iterator with it all the time. The perf shows a little bit more expensive but mostly on par with range iterator query.
15. **title:** SAMZA-1705: Switch to use snapshot in iterable impl of RocksDb
body: We should use `rocksDb.snapshot()` method to keep the snapshot and creates a new iterator with it all the time. The perf shows a little bit more expensive but mostly on par with range iterator query.

github_pulls_comments:

github_pulls_reviews:

1. Maybe add a new DBOperation here for SNAPSHOT.
2. Thanks for the catch. Let me update it.
3. How do u like this re-word? "Returns a snapshot of this store for a sorted range of entries specified by `from` to `to`. The snapshot is immutable - ie., any mutations to the store are not reflected in the snapshot after it is created"
4. **body:** sounds good. Let me update the javadoc with this.
label: documentation
5. Consider throwing an `IllegalStateException` when calling `iterator()` on a snapshot that is already closed. Also do this for the "InMemory" store case so that the behavior is consistent?
6. **body:** Nit: Add javadocs for the `KeyValueSnapshot`. "An immutable view of the {@link KeyValueStore} at a point-in-time"
label: documentation
7. **body:** nit: I re-worded the docs a little bit. consider this? "Closes this snapshot releasing any associated resources. Once a snapshot is closed, no new iterators can be created for it"
label: code-design
8. Nit: Add javadocs here. "Creates a new iterator for this snapshot. The iterator MUST be closed after its execution by invoking {@link KeyValueIterator#close}."
9. done.
10. fixed.
11. done.
12. Can we add a similar comment to the iterator method that this snapshot MUST be closed after use.
13. done.

jira_issues:

1. **summary:** Switch to use snapshot in iterable impl of RocksDb
description:
2. **summary:** Switch to use snapshot in iterable impl of RocksDb
description:

jira_issues_comments:

1. GitHub user xinyuiscool opened a pull request: <https://github.com/apache/samza/pull/510> SAMZA-1705: Switch to use snapshot in iterable impl of RocksDb We should use rocksDb.snapshot() method to keep the snapshot and creates a new iterator with it all the time. The perf shows a little bit more expensive but mostly on par with range iterator query. You can merge this pull request into a Git repository by running: \$ git pull <https://github.com/xinyuiscool/samza> SAMZA-1705 Alternatively you can review and apply these changes as the patch at: <https://github.com/apache/samza/pull/510.patch> To close this pull request, make a commit to your master/trunk branch with (at least) the following in the commit message: This closes #510 ----- commit 78841d03aa280af17b301bf2d15e877074db0aac Author: xinyuiscool <xiliu@...> Date: 2018-05-08T20:11:14Z SAMZA-1705: Switch to use snapshot in iterable impl of RocksDb -----
2. Github user asfgit closed the pull request at: <https://github.com/apache/samza/pull/510>