

Item 156

git_comments:

git_commits:

1. **summary:** commit 468ca72f2f025906cbf114f2ad62ea5e4361c5fb
message: commit 468ca72f2f025906cbf114f2ad62ea5e4361c5fb Author: Devaraj Das <ddas@yahoo-inc.com> Date: Mon Jun 28 12:05:46 2010 -0700 MAPREDUCE:1900 from <https://issues.apache.org/jira/secure/attachment/12448230/mapred-fs-close.patch> +++ b/YAHOO-CHANGES.txt + MAPREDUCE-1900. Makes JobTracker/TaskTracker close filesystems, created + on behalf of users, when they are no longer needed. (ddas) + git-svn-id: <https://svn.apache.org/repos/asf/hadoop/common/branches/branch-0.20-security-patches@107756413f79535-47bb-0310-9956-ffa450edef68>

github_issues:

github_issues_comments:

github_pulls:

github_pulls_comments:

github_pulls_reviews:

jira_issues:

1. **summary:** MapReduce daemons should close FileSystems that are not needed anymore
description: Related to HADOOP-6843, this jira is to make MapReduce behave better with respect to closing FileSystems when they are not needed anymore.
2. **summary:** MapReduce daemons should close FileSystems that are not needed anymore
description: Related to HADOOP-6843, this jira is to make MapReduce behave better with respect to closing FileSystems when they are not needed anymore.
3. **summary:** MapReduce daemons should close FileSystems that are not needed anymore
description: Related to HADOOP-6843, this jira is to make MapReduce behave better with respect to closing FileSystems when they are not needed anymore.
4. **summary:** MapReduce daemons should close FileSystems that are not needed anymore
description: Related to HADOOP-6843, this jira is to make MapReduce behave better with respect to closing FileSystems when they are not needed anymore.
5. **summary:** MapReduce daemons should close FileSystems that are not needed anymore
description: Related to HADOOP-6843, this jira is to make MapReduce behave better with respect to closing FileSystems when they are not needed anymore.
6. **summary:** MapReduce daemons should close FileSystems that are not needed anymore
description: Related to HADOOP-6843, this jira is to make MapReduce behave better with respect to closing FileSystems when they are not needed anymore.
7. **summary:** MapReduce daemons should close FileSystems that are not needed anymore
description: Related to HADOOP-6843, this jira is to make MapReduce behave better with respect to closing FileSystems when they are not needed anymore.
label: test
8. **summary:** MapReduce daemons should close FileSystems that are not needed anymore
description: Related to HADOOP-6843, this jira is to make MapReduce behave better with respect to closing FileSystems when they are not needed anymore.
9. **summary:** MapReduce daemons should close FileSystems that are not needed anymore
description: Related to HADOOP-6843, this jira is to make MapReduce behave better with respect to closing FileSystems when they are not needed anymore.
10. **summary:** MapReduce daemons should close FileSystems that are not needed anymore
description: Related to HADOOP-6843, this jira is to make MapReduce behave better with respect to closing FileSystems when they are not needed anymore.
11. **summary:** MapReduce daemons should close FileSystems that are not needed anymore

description: Related to HADOOP-6843, this jira is to make MapReduce behave better with respect to closing FileSystems when they are not needed anymore.

jira_issues_comments:

1. Patch for Y20S. The patch basically closes the filesystems, created during the job execution, at the end of the job.
2. //sigh. Had we implemented ref-counted fs cache (HADOOP-4665) and call close() every time a local context finishes using an fs instance, we are unlikely hit this problem.
3. I mean to say HADOOP-4655 instead of HADOOP-4665.
4. Missed one case. This patch fixes it.
5. a port for trunk.
6. Here is the test-patch result. Also ran ant test and test-contrib. TestSimulatorDeterministicReplay failed, but it also fails without the patch. Other tests are passing. [exec] -1 overall. [exec] [exec] +1 @author. The patch does not contain any @author tags. [exec] [exec] -1 tests included. The patch doesn't appear to include any new or modified tests. [exec] Please justify why no new tests are needed for this patch. [exec] Also please list what manual steps were performed to verify this patch. [exec] [exec] +1 javadoc. The javadoc tool did not generate any warning messages. [exec] [exec] +1 javac. The applied patch does not increase the total number of javac compiler warnings. [exec] [exec] +1 findbugs. The patch does not introduce any new Findbugs warnings. [exec] [exec] +1 release audit. The applied patch does not increase the total number of release audit warnings.
7. **body:** +1 patch looks good. "Please justify why no new tests are needed for this patch. Also please list what manual steps were performed to verify this patch."
label: test
8. The patch was tested manually on Yahoo clusters. I just committed this. Thanks, Kan for the trunk patch.
9. This patch breaks hadoop for me. Can we reconsider it? It is not safe to call `FileSystem.closeAllForUGI(UserGroupInformation.getCurrentUser());` in order to close one filesystem connection. See MAPREDUCE-2060 It causes the parent JobTracker's FileSystem object to be closed as well, making it impossible (for me) to successfully submit a job.
10. Integrated in Hadoop-Mapreduce-trunk-Commit #523 (See <https://hudson.apache.org/hudson/job/Hadoop-Mapreduce-trunk-Commit/523/>)