

**git\_comments:**

1. \* \* Reopen a new input stream with the specified position \* @param pos the position to reopen a new stream \* @throws IOException
2. every read should have thrown 1 IOException except for the first read  $144 = 128 - 1 + (128 / 8)$

**git\_commits:**

1. **summary:** HADOOP-11730. Regression: s3n read failure recovery broken. (Takenori Sato via stevel)  
**message:** HADOOP-11730. Regression: s3n read failure recovery broken. (Takenori Sato via stevel) (cherry picked from commit a6a5d1d6b5ee76c829ba7b54a4ad619f7b986681)

**github\_issues:****github\_issues\_comments:****github\_pulls:****github\_pulls\_comments:****github\_pulls\_reviews:****jira\_issues:**

1. **summary:** Regression: s3n read failure recovery broken  
**description:** s3n attempts to read again when it encounters IOException during read. But the current logic does not reopen the connection, thus, it ends up with no-op, and committing the wrong(truncated) output. Here's a stack trace as an example. {quote} 2015-03-13 20:17:24,835 [TezChild] INFO org.apache.pig.backend.hadoop.executionengine.tez.runtime.PigProcessor - Starting output org.apache.tez.mapreduce.output.MROutput@52008dbd to vertex scope-12 2015-03-13 20:17:24,866 [TezChild] DEBUG org.jets3t.service.impl.rest.httpclient.HttpMethodReleaseInputStream - Released HttpMethod as its response data stream threw an exception org.apache.http.ConnectionClosedException: Premature end of Content-Length delimited message body (expected: 296587138; received: 155648 at org.apache.http.impl.io.ContentLengthInputStream.read(ContentLengthInputStream.java:184) at org.apache.http.conn.EofSensorInputStream.read(EofSensorInputStream.java:138) at org.jets3t.service.io.InterruptableInputStream.read(InterruptableInputStream.java:78) at org.jets3t.service.impl.rest.httpclient.HttpMethodReleaseInputStream.read(HttpMethodReleaseInputStream.java:146) at org.apache.hadoop.fs.s3native.NativeS3FileSystem\$NativeS3FsInputStream.read(NativeS3FileSystem.java:145) at java.io.BufferedInputStream.read1(BufferedInputStream.java:273) at java.io.BufferedInputStream.read(BufferedInputStream.java:334) at java.io.DataInputStream.read(DataInputStream.java:100) at org.apache.hadoop.util.LineReader.fillBuffer(LineReader.java:180) at org.apache.hadoop.util.LineReader.readDefaultLine(LineReader.java:216) at org.apache.hadoop.util.LineReader.readLine(LineReader.java:174) at org.apache.hadoop.mapreduce.lib.input.LineRecordReader.nextKeyValue(LineRecordReader.java:185) at org.apache.pig.builtin.PigStorage.getNext(PigStorage.java:259) at org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.PigRecordReader.nextKeyValue(PigRecordReader.java:204) at org.apache.tez.mapreduce.lib.MRReaderMapReduce.next(MRReaderMapReduce.java:116) at org.apache.pig.backend.hadoop.executionengine.tez.plan.operator.POSimpleTezLoad.getNextTuple(POSimpleTezLoad.java:106) at org.apache.pig.backend.hadoop.executionengine.physicalLayer.PhysicalOperator.processInput(PhysicalOperator.java:307) at org.apache.pig.backend.hadoop.executionengine.physicalLayer.relationalOperators.POForEach.getNextTuple(POForEach.java:246) at org.apache.pig.backend.hadoop.executionengine.physicalLayer.PhysicalOperator.processInput(PhysicalOperator.java:307) at org.apache.pig.backend.hadoop.executionengine.physicalLayer.relationalOperators.POFilter.getNextTuple(POFilter.java:91) at org.apache.pig.backend.hadoop.executionengine.physicalLayer.PhysicalOperator.processInput(PhysicalOperator.java:307) at org.apache.pig.backend.hadoop.executionengine.tez.plan.operator.POStoreTez.getNextTuple(POStoreTez.java:117) at org.apache.pig.backend.hadoop.executionengine.tez.runtime.PigProcessor.runPipeline(PigProcessor.java:313) at org.apache.pig.backend.hadoop.executionengine.tez.runtime.PigProcessor.run(PigProcessor.java:192) at org.apache.tez.runtime.LogicalIOProcessorRuntimeTask.run(LogicalIOProcessorRuntimeTask.java:324) at org.apache.tez.runtime.task.TezTaskRunner\$TaskRunnerCallable\$1.run(TezTaskRunner.java:176) at org.apache.tez.runtime.task.TezTaskRunner\$TaskRunnerCallable\$1.run(TezTaskRunner.java:168) at java.security.AccessController.doPrivileged(Native Method) at javax.security.auth.Subject.doAs(Subject.java:415) at org.apache.hadoop.security.UserGroupInformation.doAs(UserGroupInformation.java:1628) at org.apache.tez.runtime.task.TezTaskRunner\$TaskRunnerCallable.call(TezTaskRunner.java:168) at org.apache.tez.runtime.task.TezTaskRunner\$TaskRunnerCallable.call(TezTaskRunner.java:163) at java.util.concurrent.FutureTask.run(FutureTask.java:262) at java.util.concurrent.ThreadPoolExecutor.runWorker(ThreadPoolExecutor.java:1145) at java.util.concurrent.ThreadPoolExecutor\$Worker.run(ThreadPoolExecutor.java:615) at java.lang.Thread.run(Thread.java:745) 2015-03-13 20:17:24,867 [TezChild] INFO org.apache.hadoop.fs.s3native.NativeS3FileSystem - Received IOException while reading 'user/hadoop/tsato/readlarge/input/cloudian-s3.log.20141119', attempting to reopen. 2015-03-13 20:17:24,867 [TezChild] DEBUG org.jets3t.service.impl.rest.httpclient.HttpMethodReleaseInputStream - Released HttpMethod as its response data stream is fully consumed 2015-03-13 20:17:24,868 [TezChild] INFO org.apache.tez.dag.app.TaskAttemptListenerImpTezDag - Commit go/no-go request from attempt\_1426245338920\_0001\_1\_00\_000004\_0 2015-03-13 20:17:24,868 [TezChild] INFO org.apache.tez.dag.app.dag.impl.TaskImpl - attempt\_1426245338920\_0001\_1\_00\_000004\_0 given a go for committing the task output. {quote} It seems this is a regression, which was introduced by the following optimizations.

<https://issues.apache.org/jira/browse/HADOOP-10589> <https://issues.apache.org/jira/browse/HADOOP-10457> Also, test cases should be reviewed so that it covers this scenario.

**label:** code-design

2. **summary:** Regression: s3n read failure recovery broken

**description:** s3n attempts to read again when it encounters IOException during read. But the current logic does not reopen the connection, thus, it ends up with no-op, and committing the wrong(truncated) output. Here's a stack trace as an example. {quote} 2015-03-13 20:17:24,835 [TezChild] INFO org.apache.pig.backend.hadoop.executionengine.tez.runtime.PigProcessor - Starting output org.apache.tez.mapreduce.output.MROutput@52008dbd to vertex scope-12 2015-03-13 20:17:24,866 [TezChild] DEBUG org.jets3t.service.impl.rest.httpclient.HttpMethodReleaseInputStream - Released HttpMethod as its response data stream threw an exception org.apache.http.ConnectionClosedException: Premature end of Content-Length delimited message body (expected: 296587138; received: 155648 at org.apache.http.impl.io.ContentLengthInputStream.read(ContentLengthInputStream.java:184) at org.apache.http.conn.EofSensorInputStream.read(EofSensorInputStream.java:138) at org.jets3t.service.io.InterruptableInputStream.read(InterruptableInputStream.java:78) at org.jets3t.service.impl.rest.httpclient.HttpMethodReleaseInputStream.read(HttpMethodReleaseInputStream.java:146) at org.apache.hadoop.fs.s3native.NativeS3FileSystem\$NativeS3FsInputStream.read(NativeS3FileSystem.java:145) at java.io.BufferedInputStream.read1(BufferedInputStream.java:273) at java.io.BufferedInputStream.read(BufferedInputStream.java:334) at java.io.DataInputStream.read(DataInputStream.java:100) at org.apache.hadoop.util.LineReader.fillBuffer(LineReader.java:180) at org.apache.hadoop.util.LineReader.readDefaultLine(LineReader.java:216) at org.apache.hadoop.util.LineReader.readLine(LineReader.java:174) at org.apache.hadoop.mapreduce.lib.input.LineRecordReader.nextKeyValue(LineRecordReader.java:185) at org.apache.pig.builtin.PigStorage.getNext(PigStorage.java:259) at org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.PigRecordReader.nextKeyValue(PigRecordReader.java:204) at org.apache.tez.mapreduce.lib.MRReaderMapReduce.next(MRReaderMapReduce.java:116) at org.apache.pig.backend.hadoop.executionengine.tez.plan.operator.POSimpleTezLoad.getNextTuple(POSimpleTezLoad.java:106) at org.apache.pig.backend.hadoop.executionengine.physicalLayer.PhysicalOperator.processInput(PhysicalOperator.java:307) at org.apache.pig.backend.hadoop.executionengine.physicalLayer.relationalOperators.POForEach.getNextTuple(POForEach.java:246) at org.apache.pig.backend.hadoop.executionengine.physicalLayer.PhysicalOperator.processInput(PhysicalOperator.java:307) at org.apache.pig.backend.hadoop.executionengine.physicalLayer.relationalOperators.POFilter.getNextTuple(POFilter.java:91) at org.apache.pig.backend.hadoop.executionengine.physicalLayer.PhysicalOperator.processInput(PhysicalOperator.java:307) at org.apache.pig.backend.hadoop.executionengine.tez.plan.operator.POStoreTez.getNextTuple(POStoreTez.java:117) at org.apache.pig.backend.hadoop.executionengine.tez.runtime.PigProcessor.runPipeline(PigProcessor.java:313) at org.apache.pig.backend.hadoop.executionengine.tez.runtime.PigProcessor.run(PigProcessor.java:192) at org.apache.tez.runtime.LogicalIOProcessorRuntimeTask.run(LogicalIOProcessorRuntimeTask.java:324) at org.apache.tez.runtime.task.TezTaskRunner\$TaskRunnerCallable\$1.run(TezTaskRunner.java:176) at org.apache.tez.runtime.task.TezTaskRunner\$TaskRunnerCallable\$1.run(TezTaskRunner.java:168) at java.security.AccessController.doPrivileged(Native Method) at javax.security.auth.Subject.doAs(Subject.java:415) at org.apache.hadoop.security.UserGroupInformation.doAs(UserGroupInformation.java:1628) at org.apache.tez.runtime.task.TezTaskRunner\$TaskRunnerCallable.call(TezTaskRunner.java:168) at org.apache.tez.runtime.task.TezTaskRunner\$TaskRunnerCallable.call(TezTaskRunner.java:163) at java.util.concurrent.FutureTask.run(FutureTask.java:262) at java.util.concurrent.ThreadPoolExecutor.runWorker(ThreadPoolExecutor.java:1145) at java.util.concurrent.ThreadPoolExecutor\$Worker.run(ThreadPoolExecutor.java:615) at java.lang.Thread.run(Thread.java:745) 2015-03-13 20:17:24,867 [TezChild] INFO org.apache.hadoop.fs.s3native.NativeS3FileSystem - Received IOException while reading 'user/hadoop/tsato/readlarge/input/cloudian-s3.log.20141119', attempting to reopen. 2015-03-13 20:17:24,867 [TezChild] DEBUG org.jets3t.service.impl.rest.httpclient.HttpMethodReleaseInputStream - Released HttpMethod as its response data stream is fully consumed 2015-03-13 20:17:24,868 [TezChild] INFO org.apache.tez.dag.app.TaskAttemptListenerImplTezDag - Commit go/no-go request from attempt\_1426245338920\_0001\_1\_00\_000004\_0 2015-03-13 20:17:24,868 [TezChild] INFO org.apache.tez.dag.app.dag.impl.TaskImpl - attempt\_1426245338920\_0001\_1\_00\_000004\_0 given a go for committing the task output. {quote} It seems this is a regression, which was introduced by the following optimizations. <https://issues.apache.org/jira/browse/HADOOP-10589> <https://issues.apache.org/jira/browse/HADOOP-10457> Also, test cases should be reviewed so that it covers this scenario.

3. **summary:** Regression: s3n read failure recovery broken

**description:** s3n attempts to read again when it encounters IOException during read. But the current logic does not reopen the connection, thus, it ends up with no-op, and committing the wrong(truncated) output. Here's a stack trace as an example. {quote} 2015-03-13 20:17:24,835 [TezChild] INFO org.apache.pig.backend.hadoop.executionengine.tez.runtime.PigProcessor - Starting output org.apache.tez.mapreduce.output.MROutput@52008dbd to vertex scope-12 2015-03-13 20:17:24,866 [TezChild] DEBUG org.jets3t.service.impl.rest.httpclient.HttpMethodReleaseInputStream - Released HttpMethod as its response data stream threw an exception org.apache.http.ConnectionClosedException: Premature end of Content-Length delimited message body (expected: 296587138; received: 155648 at org.apache.http.impl.io.ContentLengthInputStream.read(ContentLengthInputStream.java:184) at org.apache.http.conn.EofSensorInputStream.read(EofSensorInputStream.java:138) at org.jets3t.service.io.InterruptableInputStream.read(InterruptableInputStream.java:78) at org.jets3t.service.impl.rest.httpclient.HttpMethodReleaseInputStream.read(HttpMethodReleaseInputStream.java:146) at org.apache.hadoop.fs.s3native.NativeS3FileSystem\$NativeS3FsInputStream.read(NativeS3FileSystem.java:145) at java.io.BufferedInputStream.read1(BufferedInputStream.java:273) at java.io.BufferedInputStream.read(BufferedInputStream.java:334) at java.io.DataInputStream.read(DataInputStream.java:100) at org.apache.hadoop.util.LineReader.fillBuffer(LineReader.java:180) at org.apache.hadoop.util.LineReader.readDefaultLine(LineReader.java:216) at org.apache.hadoop.util.LineReader.readLine(LineReader.java:174) at org.apache.hadoop.mapreduce.lib.input.LineRecordReader.nextKeyValue(LineRecordReader.java:185) at org.apache.pig.builtin.PigStorage.getNext(PigStorage.java:259) at

org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.PigRecordReader.nextKeyValue(PigRecordReader.java:204) at  
 org.apache.tez.mapreduce.lib.MRReaderMapReduce.next(MRReaderMapReduce.java:116) at  
 org.apache.pig.backend.hadoop.executionengine.tez.plan.operator.POSimpleTezLoad.getNextTuple(POSimpleTezLoad.java:106)  
 at org.apache.pig.backend.hadoop.executionengine.physicalLayer.PhysicalOperator.processInput(PhysicalOperator.java:307) at  
 org.apache.pig.backend.hadoop.executionengine.physicalLayer.relationalOperators.POForEach.getNextTuple(POForEach.java:246)  
 at org.apache.pig.backend.hadoop.executionengine.physicalLayer.PhysicalOperator.processInput(PhysicalOperator.java:307) at  
 org.apache.pig.backend.hadoop.executionengine.physicalLayer.relationalOperators.POFilter.getNextTuple(POFilter.java:91) at  
 org.apache.pig.backend.hadoop.executionengine.physicalLayer.PhysicalOperator.processInput(PhysicalOperator.java:307) at  
 org.apache.pig.backend.hadoop.executionengine.tez.plan.operator.POSoreTez.getNextTuple(POSoreTez.java:117) at  
 org.apache.pig.backend.hadoop.executionengine.tez.runtime.PigProcessor.runPipeline(PigProcessor.java:313) at  
 org.apache.pig.backend.hadoop.executionengine.tez.runtime.PigProcessor.run(PigProcessor.java:192) at  
 org.apache.tez.runtime.LogicalIOProcessorRuntimeTask.run(LogicalIOProcessorRuntimeTask.java:324) at  
 org.apache.tez.runtime.task.TezTaskRunner\$TaskRunnerCallable\$1.run(TezTaskRunner.java:176) at  
 org.apache.tez.runtime.task.TezTaskRunner\$TaskRunnerCallable\$1.run(TezTaskRunner.java:168) at  
 java.security.AccessController.doPrivileged(Native Method) at javax.security.auth.Subject.doAs(Subject.java:415) at  
 org.apache.hadoop.security.UserGroupInformation.doAs(UserGroupInformation.java:1628) at  
 org.apache.tez.runtime.task.TezTaskRunner\$TaskRunnerCallable.call(TezTaskRunner.java:168) at  
 org.apache.tez.runtime.task.TezTaskRunner\$TaskRunnerCallable.call(TezTaskRunner.java:163) at  
 java.util.concurrent.FutureTask.run(FutureTask.java:262) at  
 java.util.concurrent.ThreadPoolExecutor.runWorker(ThreadPoolExecutor.java:1145) at  
 java.util.concurrent.ThreadPoolExecutor\$Worker.run(ThreadPoolExecutor.java:615) at java.lang.Thread.run(Thread.java:745)  
 2015-03-13 20:17:24,867 [TezChild] INFO org.apache.hadoop.fs.s3native.NativeS3FileSystem - Received IOException while  
 reading 'user/hadoop/tsato/readlarge/input/cloudian-s3.log.20141119', attempting to reopen. 2015-03-13 20:17:24,867 [TezChild]  
 DEBUG org.jets3t.service.impl.rest.httpclient.HttpMethodReleaseInputStream - Released HttpMethod as its response data stream  
 is fully consumed 2015-03-13 20:17:24,868 [TezChild] INFO org.apache.tez.dag.app.TaskAttemptListenerImpTezDag - Commit  
 go/no-go request from attempt\_1426245338920\_0001\_1\_00\_000004\_0 2015-03-13 20:17:24,868 [TezChild] INFO  
 org.apache.tez.dag.app.dag.impl.TaskImpl - attempt\_1426245338920\_0001\_1\_00\_000004\_0 given a go for committing the task  
 output. {quote} It seems this is a regression, which was introduced by the following optimizations.  
<https://issues.apache.org/jira/browse/HADOOP-10589> <https://issues.apache.org/jira/browse/HADOOP-10457> Also, test cases  
 should be reviewed so that it covers this scenario.

#### 4. **summary:** Regression: s3n read failure recovery broken

**description:** s3n attempts to read again when it encounters IOException during read. But the current logic does not reopen the  
 connection, thus, it ends up with no-op, and committing the wrong(truncated) output. Here's a stack trace as an example. {quote}  
 2015-03-13 20:17:24,835 [TezChild] INFO org.apache.pig.backend.hadoop.executionengine.tez.runtime.PigProcessor - Starting  
 output org.apache.tez.mapreduce.output.MROutput@52008dbd to vertex scope-12 2015-03-13 20:17:24,866 [TezChild] DEBUG  
 org.jets3t.service.impl.rest.httpclient.HttpMethodReleaseInputStream - Released HttpMethod as its response data stream threw an  
 exception org.apache.http.ConnectionClosedException: Premature end of Content-Length delimited message body (expected:  
 296587138; received: 155648 at org.apache.http.impl.io.ContentLengthInputStream.read(ContentLengthInputStream.java:184) at  
 org.apache.http.conn.EofSensorInputStream.read(EofSensorInputStream.java:138) at  
 org.jets3t.service.io.InterruptableInputStream.read(InterruptableInputStream.java:78) at  
 org.jets3t.service.impl.rest.httpclient.HttpMethodReleaseInputStream.read(HttpMethodReleaseInputStream.java:146) at  
 org.apache.hadoop.fs.s3native.NativeS3FileSystem\$NativeS3FsInputStream.read(NativeS3FileSystem.java:145) at  
 java.io.BufferedInputStream.read1(BufferedInputStream.java:273) at  
 java.io.BufferedInputStream.read(BufferedInputStream.java:334) at java.io.DataInputStream.read(DataInputStream.java:100) at  
 org.apache.hadoop.util.LineReader.fillBuffer(LineReader.java:180) at  
 org.apache.hadoop.util.LineReader.readDefaultLine(LineReader.java:216) at  
 org.apache.hadoop.util.LineReader.readLine(LineReader.java:174) at  
 org.apache.hadoop.mapreduce.lib.input.LineRecordReader.nextKeyValue(LineRecordReader.java:185) at  
 org.apache.pig.builtin.PigStorage.getNext(PigStorage.java:259) at  
 org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.PigRecordReader.nextKeyValue(PigRecordReader.java:204) at  
 org.apache.tez.mapreduce.lib.MRReaderMapReduce.next(MRReaderMapReduce.java:116) at  
 org.apache.pig.backend.hadoop.executionengine.tez.plan.operator.POSimpleTezLoad.getNextTuple(POSimpleTezLoad.java:106)  
 at org.apache.pig.backend.hadoop.executionengine.physicalLayer.PhysicalOperator.processInput(PhysicalOperator.java:307) at  
 org.apache.pig.backend.hadoop.executionengine.physicalLayer.relationalOperators.POForEach.getNextTuple(POForEach.java:246)  
 at org.apache.pig.backend.hadoop.executionengine.physicalLayer.PhysicalOperator.processInput(PhysicalOperator.java:307) at  
 org.apache.pig.backend.hadoop.executionengine.physicalLayer.relationalOperators.POFilter.getNextTuple(POFilter.java:91) at  
 org.apache.pig.backend.hadoop.executionengine.physicalLayer.PhysicalOperator.processInput(PhysicalOperator.java:307) at  
 org.apache.pig.backend.hadoop.executionengine.tez.plan.operator.POSoreTez.getNextTuple(POSoreTez.java:117) at  
 org.apache.pig.backend.hadoop.executionengine.tez.runtime.PigProcessor.runPipeline(PigProcessor.java:313) at  
 org.apache.pig.backend.hadoop.executionengine.tez.runtime.PigProcessor.run(PigProcessor.java:192) at  
 org.apache.pig.backend.hadoop.executionengine.tez.runtime.PigProcessor.run(PigProcessor.java:192) at  
 org.apache.tez.runtime.LogicalIOProcessorRuntimeTask.run(LogicalIOProcessorRuntimeTask.java:324) at  
 org.apache.tez.runtime.task.TezTaskRunner\$TaskRunnerCallable\$1.run(TezTaskRunner.java:176) at  
 org.apache.tez.runtime.task.TezTaskRunner\$TaskRunnerCallable\$1.run(TezTaskRunner.java:168) at  
 java.security.AccessController.doPrivileged(Native Method) at javax.security.auth.Subject.doAs(Subject.java:415) at  
 org.apache.hadoop.security.UserGroupInformation.doAs(UserGroupInformation.java:1628) at  
 org.apache.tez.runtime.task.TezTaskRunner\$TaskRunnerCallable.call(TezTaskRunner.java:168) at  
 org.apache.tez.runtime.task.TezTaskRunner\$TaskRunnerCallable.call(TezTaskRunner.java:163) at  
 java.util.concurrent.FutureTask.run(FutureTask.java:262) at  
 java.util.concurrent.ThreadPoolExecutor.runWorker(ThreadPoolExecutor.java:1145) at  
 java.util.concurrent.ThreadPoolExecutor\$Worker.run(ThreadPoolExecutor.java:615) at java.lang.Thread.run(Thread.java:745)  
 2015-03-13 20:17:24,867 [TezChild] INFO org.apache.hadoop.fs.s3native.NativeS3FileSystem - Received IOException while

reading 'user/hadoop/tsato/readlarge/input/cloudian-s3.log.20141119', attempting to reopen. 2015-03-13 20:17:24,867 [TezChild] DEBUG org.jets3t.service.impl.rest.httpclient.HttpMethodReleaseInputStream - Released HttpMethod as its response data stream is fully consumed 2015-03-13 20:17:24,868 [TezChild] INFO org.apache.tez.dag.app.TaskAttemptListenerImpTezDag - Commit go/no-go request from attempt\_1426245338920\_0001\_1\_00\_000004\_0 2015-03-13 20:17:24,868 [TezChild] INFO org.apache.tez.dag.app.dag.impl.TaskImpl - attempt\_1426245338920\_0001\_1\_00\_000004\_0 given a go for committing the task output. {quote} It seems this is a regression, which was introduced by the following optimizations. <https://issues.apache.org/jira/browse/HADOOP-10589> <https://issues.apache.org/jira/browse/HADOOP-10457> Also, test cases should be reviewed so that it covers this scenario.

5. **summary:** Regression: s3n read failure recovery broken

**description:** s3n attempts to read again when it encounters IOException during read. But the current logic does not reopen the connection, thus, it ends up with no-op, and committing the wrong(truncated) output. Here's a stack trace as an example. {quote} 2015-03-13 20:17:24,835 [TezChild] INFO org.apache.pig.backend.hadoop.executionengine.tez.runtime.PigProcessor - Starting output org.apache.tez.mapreduce.output.MROutput@52008dbd to vertex scope-12 2015-03-13 20:17:24,866 [TezChild] DEBUG org.jets3t.service.impl.rest.httpclient.HttpMethodReleaseInputStream - Released HttpMethod as its response data stream threw an exception org.apache.http.ConnectionClosedException: Premature end of Content-Length delimited message body (expected: 296587138; received: 155648 at org.apache.http.impl.io.ContentLengthInputStream.read(ContentLengthInputStream.java:184) at org.apache.http.conn.EofSensorInputStream.read(EofSensorInputStream.java:138) at org.jets3t.service.io.InterruptableInputStream.read(InterruptableInputStream.java:78) at org.jets3t.service.impl.rest.httpclient.HttpMethodReleaseInputStream.read(HttpMethodReleaseInputStream.java:146) at org.apache.hadoop.fs.s3native.NativeS3FileSystem\$NativeS3FsInputStream.read(NativeS3FileSystem.java:145) at java.io.BufferedInputStream.read1(BufferedInputStream.java:273) at java.io.BufferedInputStream.read(BufferedInputStream.java:334) at java.io.DataInputStream.read(DataInputStream.java:100) at org.apache.hadoop.util.LineReader.fillBuffer(LineReader.java:180) at org.apache.hadoop.util.LineReader.readDefaultLine(LineReader.java:216) at org.apache.hadoop.util.LineReader.readLine(LineReader.java:174) at org.apache.hadoop.mapreduce.lib.input.LineRecordReader.nextKeyValue(LineRecordReader.java:185) at org.apache.pig.builtin.PigStorage.getNext(PigStorage.java:259) at org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.PigRecordReader.nextKeyValue(PigRecordReader.java:204) at org.apache.tez.mapreduce.lib.MRReaderMapReduce.next(MRReaderMapReduce.java:116) at org.apache.pig.backend.hadoop.executionengine.tez.plan.operator.POSimpleTezLoad.getNextTuple(POSimpleTezLoad.java:106) at org.apache.pig.backend.hadoop.executionengine.physicalLayer.PhysicalOperator.processInput(PhysicalOperator.java:307) at org.apache.pig.backend.hadoop.executionengine.physicalLayer.relationalOperators.POForEach.getNextTuple(POForEach.java:246) at org.apache.pig.backend.hadoop.executionengine.physicalLayer.PhysicalOperator.processInput(PhysicalOperator.java:307) at org.apache.pig.backend.hadoop.executionengine.physicalLayer.relationalOperators.POFilter.getNextTuple(POFilter.java:91) at org.apache.pig.backend.hadoop.executionengine.physicalLayer.PhysicalOperator.processInput(PhysicalOperator.java:307) at org.apache.pig.backend.hadoop.executionengine.tez.plan.operator.POStoreTez.getNextTuple(POStoreTez.java:117) at org.apache.pig.backend.hadoop.executionengine.tez.runtime.PigProcessor.runPipeline(PigProcessor.java:313) at org.apache.pig.backend.hadoop.executionengine.tez.runtime.PigProcessor.run(PigProcessor.java:192) at org.apache.tez.runtime.LogicalIOProcessorRuntimeTask.run(LogicalIOProcessorRuntimeTask.java:324) at org.apache.tez.runtime.task.TezTaskRunner\$TaskRunnerCallable\$1.run(TezTaskRunner.java:176) at org.apache.tez.runtime.task.TezTaskRunner\$TaskRunnerCallable\$1.run(TezTaskRunner.java:168) at java.security.AccessController.doPrivileged(Native Method) at javax.security.auth.Subject.doAs(Subject.java:415) at org.apache.hadoop.security.UserGroupInformation.doAs(UserGroupInformation.java:1628) at org.apache.tez.runtime.task.TezTaskRunner\$TaskRunnerCallable.call(TezTaskRunner.java:168) at org.apache.tez.runtime.task.TezTaskRunner\$TaskRunnerCallable.call(TezTaskRunner.java:163) at java.util.concurrent.FutureTask.run(FutureTask.java:262) at java.util.concurrent.ThreadPoolExecutor.runWorker(ThreadPoolExecutor.java:1145) at java.util.concurrent.ThreadPoolExecutor\$Worker.run(ThreadPoolExecutor.java:615) at java.lang.Thread.run(Thread.java:745) 2015-03-13 20:17:24,867 [TezChild] INFO org.apache.hadoop.fs.s3native.NativeS3FileSystem - Received IOException while reading 'user/hadoop/tsato/readlarge/input/cloudian-s3.log.20141119', attempting to reopen. 2015-03-13 20:17:24,867 [TezChild] DEBUG org.jets3t.service.impl.rest.httpclient.HttpMethodReleaseInputStream - Released HttpMethod as its response data stream is fully consumed 2015-03-13 20:17:24,868 [TezChild] INFO org.apache.tez.dag.app.TaskAttemptListenerImpTezDag - Commit go/no-go request from attempt\_1426245338920\_0001\_1\_00\_000004\_0 2015-03-13 20:17:24,868 [TezChild] INFO org.apache.tez.dag.app.dag.impl.TaskImpl - attempt\_1426245338920\_0001\_1\_00\_000004\_0 given a go for committing the task output. {quote} It seems this is a regression, which was introduced by the following optimizations. <https://issues.apache.org/jira/browse/HADOOP-10589> <https://issues.apache.org/jira/browse/HADOOP-10457> Also, test cases should be reviewed so that it covers this scenario.

**label:** test

6. **summary:** Regression: s3n read failure recovery broken

**description:** s3n attempts to read again when it encounters IOException during read. But the current logic does not reopen the connection, thus, it ends up with no-op, and committing the wrong(truncated) output. Here's a stack trace as an example. {quote} 2015-03-13 20:17:24,835 [TezChild] INFO org.apache.pig.backend.hadoop.executionengine.tez.runtime.PigProcessor - Starting output org.apache.tez.mapreduce.output.MROutput@52008dbd to vertex scope-12 2015-03-13 20:17:24,866 [TezChild] DEBUG org.jets3t.service.impl.rest.httpclient.HttpMethodReleaseInputStream - Released HttpMethod as its response data stream threw an exception org.apache.http.ConnectionClosedException: Premature end of Content-Length delimited message body (expected: 296587138; received: 155648 at org.apache.http.impl.io.ContentLengthInputStream.read(ContentLengthInputStream.java:184) at org.apache.http.conn.EofSensorInputStream.read(EofSensorInputStream.java:138) at org.jets3t.service.io.InterruptableInputStream.read(InterruptableInputStream.java:78) at org.jets3t.service.impl.rest.httpclient.HttpMethodReleaseInputStream.read(HttpMethodReleaseInputStream.java:146) at org.apache.hadoop.fs.s3native.NativeS3FileSystem\$NativeS3FsInputStream.read(NativeS3FileSystem.java:145) at java.io.BufferedInputStream.read1(BufferedInputStream.java:273) at

java.io.BufferedInputStream.read(BufferedInputStream.java:334) at java.io.DataInputStream.read(DataInputStream.java:100) at  
org.apache.hadoop.util.LineReader.fillBuffer(LineReader.java:180) at  
org.apache.hadoop.util.LineReader.readDefaultLine(LineReader.java:216) at  
org.apache.hadoop.util.LineReader.readLine(LineReader.java:174) at  
org.apache.hadoop.mapreduce.lib.input.LineRecordReader.nextKeyValue(LineRecordReader.java:185) at  
org.apache.pig.builtin.PigStorage.getNext(PigStorage.java:259) at  
org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.PigRecordReader.nextKeyValue(PigRecordReader.java:204) at  
org.apache.tez.mapreduce.lib.MRReaderMapReduce.next(MRReaderMapReduce.java:116) at  
org.apache.pig.backend.hadoop.executionengine.tez.plan.operator.POSimpleTezLoad.getNextTuple(POSimpleTezLoad.java:106)  
at org.apache.pig.backend.hadoop.executionengine.physicalLayer.PhysicalOperator.processInput(PhysicalOperator.java:307) at  
org.apache.pig.backend.hadoop.executionengine.physicalLayer.relationalOperators.POForEach.getNextTuple(POForEach.java:246)  
at org.apache.pig.backend.hadoop.executionengine.physicalLayer.PhysicalOperator.processInput(PhysicalOperator.java:307) at  
org.apache.pig.backend.hadoop.executionengine.physicalLayer.relationalOperators.POFilter.getNextTuple(POFilter.java:91) at  
org.apache.pig.backend.hadoop.executionengine.physicalLayer.PhysicalOperator.processInput(PhysicalOperator.java:307) at  
org.apache.pig.backend.hadoop.executionengine.tez.plan.operator.POStoreTez.getNextTuple(POStoreTez.java:117) at  
org.apache.pig.backend.hadoop.executionengine.tez.runtime.PigProcessor.runPipeline(PigProcessor.java:313) at  
org.apache.pig.backend.hadoop.executionengine.tez.runtime.PigProcessor.run(PigProcessor.java:192) at  
org.apache.tez.runtime.LogicalIOProcessorRuntimeTask.run(LogicalIOProcessorRuntimeTask.java:324) at  
org.apache.tez.runtime.task.TezTaskRunner\$TaskRunnerCallable\$1.run(TezTaskRunner.java:176) at  
org.apache.tez.runtime.task.TezTaskRunner\$TaskRunnerCallable\$1.run(TezTaskRunner.java:168) at  
java.security.AccessController.doPrivileged(Native Method) at javax.security.auth.Subject.doAs(Subject.java:415) at  
org.apache.hadoop.security.UserGroupInformation.doAs(UserGroupInformation.java:1628) at  
org.apache.tez.runtime.task.TezTaskRunner\$TaskRunnerCallable.call(TezTaskRunner.java:168) at  
org.apache.tez.runtime.task.TezTaskRunner\$TaskRunnerCallable.call(TezTaskRunner.java:163) at  
java.util.concurrent.FutureTask.run(FutureTask.java:262) at  
java.util.concurrent.ThreadPoolExecutor.runWorker(ThreadPoolExecutor.java:1145) at  
java.util.concurrent.ThreadPoolExecutor\$Worker.run(ThreadPoolExecutor.java:615) at java.lang.Thread.run(Thread.java:745)  
2015-03-13 20:17:24,867 [TezChild] INFO org.apache.hadoop.fs.s3native.NativeS3FileSystem - Received IOException while  
reading 'user/hadoop/tsato/readlarge/input/cloudian-s3.log.20141119', attempting to reopen. 2015-03-13 20:17:24,867 [TezChild]  
DEBUG org.jets3t.service.impl.rest.httpclient.HttpMethodReleaseInputStream - Released HttpMethod as its response data stream  
is fully consumed 2015-03-13 20:17:24,868 [TezChild] INFO org.apache.tez.dag.app.TaskAttemptListenerImpTezDag - Commit  
go/no-go request from attempt\_1426245338920\_0001\_1\_00\_000004\_0 2015-03-13 20:17:24,868 [TezChild] INFO  
org.apache.tez.dag.app.dag.impl.TaskImpl - attempt\_1426245338920\_0001\_1\_00\_000004\_0 given a go for committing the task  
output. {quote} It seems this is a regression, which was introduced by the following optimizations.  
<https://issues.apache.org/jira/browse/HADOOP-10589> <https://issues.apache.org/jira/browse/HADOOP-10457> Also, test cases  
should be reviewed so that it covers this scenario.

7. **summary:** Regression: s3n read failure recovery broken

**description:** s3n attempts to read again when it encounters IOException during read. But the current logic does not reopen the  
connection, thus, it ends up with no-op, and committing the wrong(truncated) output. Here's a stack trace as an example. {quote}  
2015-03-13 20:17:24,835 [TezChild] INFO org.apache.pig.backend.hadoop.executionengine.tez.runtime.PigProcessor - Starting  
output org.apache.tez.mapreduce.output.MROutput@52008dbd to vertex scope-12 2015-03-13 20:17:24,866 [TezChild] DEBUG  
org.jets3t.service.impl.rest.httpclient.HttpMethodReleaseInputStream - Released HttpMethod as its response data stream threw an  
exception org.apache.http.ConnectionClosedException: Premature end of Content-Length delimited message body (expected:  
296587138; received: 155648 at org.apache.http.impl.io.ContentLengthInputStream.read(ContentLengthInputStream.java:184) at  
org.apache.http.conn.EofSensorInputStream.read(EofSensorInputStream.java:138) at  
org.jets3t.service.io.InterruptableInputStream.read(InterruptableInputStream.java:78) at  
org.jets3t.service.impl.rest.httpclient.HttpMethodReleaseInputStream.read(HttpMethodReleaseInputStream.java:146) at  
org.apache.hadoop.fs.s3native.NativeS3FileSystem\$NativeS3FsInputStream.read(NativeS3FileSystem.java:145) at  
java.io.BufferedInputStream.read1(BufferedInputStream.java:273) at  
java.io.BufferedInputStream.read(BufferedInputStream.java:334) at java.io.DataInputStream.read(DataInputStream.java:100) at  
org.apache.hadoop.util.LineReader.fillBuffer(LineReader.java:180) at  
org.apache.hadoop.util.LineReader.readDefaultLine(LineReader.java:216) at  
org.apache.hadoop.util.LineReader.readLine(LineReader.java:174) at  
org.apache.hadoop.mapreduce.lib.input.LineRecordReader.nextKeyValue(LineRecordReader.java:185) at  
org.apache.pig.builtin.PigStorage.getNext(PigStorage.java:259) at  
org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.PigRecordReader.nextKeyValue(PigRecordReader.java:204) at  
org.apache.tez.mapreduce.lib.MRReaderMapReduce.next(MRReaderMapReduce.java:116) at  
org.apache.pig.backend.hadoop.executionengine.tez.plan.operator.POSimpleTezLoad.getNextTuple(POSimpleTezLoad.java:106)  
at org.apache.pig.backend.hadoop.executionengine.physicalLayer.PhysicalOperator.processInput(PhysicalOperator.java:307) at  
org.apache.pig.backend.hadoop.executionengine.physicalLayer.relationalOperators.POForEach.getNextTuple(POForEach.java:246)  
at org.apache.pig.backend.hadoop.executionengine.physicalLayer.PhysicalOperator.processInput(PhysicalOperator.java:307) at  
org.apache.pig.backend.hadoop.executionengine.physicalLayer.relationalOperators.POFilter.getNextTuple(POFilter.java:91) at  
org.apache.pig.backend.hadoop.executionengine.physicalLayer.PhysicalOperator.processInput(PhysicalOperator.java:307) at  
org.apache.pig.backend.hadoop.executionengine.tez.plan.operator.POStoreTez.getNextTuple(POStoreTez.java:117) at  
org.apache.pig.backend.hadoop.executionengine.tez.runtime.PigProcessor.runPipeline(PigProcessor.java:313) at  
org.apache.pig.backend.hadoop.executionengine.tez.runtime.PigProcessor.run(PigProcessor.java:192) at  
org.apache.tez.runtime.LogicalIOProcessorRuntimeTask.run(LogicalIOProcessorRuntimeTask.java:324) at  
org.apache.tez.runtime.task.TezTaskRunner\$TaskRunnerCallable\$1.run(TezTaskRunner.java:176) at  
org.apache.tez.runtime.task.TezTaskRunner\$TaskRunnerCallable\$1.run(TezTaskRunner.java:168) at  
java.security.AccessController.doPrivileged(Native Method) at javax.security.auth.Subject.doAs(Subject.java:415) at  
org.apache.hadoop.security.UserGroupInformation.doAs(UserGroupInformation.java:1628) at

org.apache.tez.runtime.task.TezTaskRunner\$TaskRunnerCallable.call(TezTaskRunner.java:168) at  
org.apache.tez.runtime.task.TezTaskRunner\$TaskRunnerCallable.call(TezTaskRunner.java:163) at  
java.util.concurrent.FutureTask.run(FutureTask.java:262) at  
java.util.concurrent.ThreadPoolExecutor.runWorker(ThreadPoolExecutor.java:1145) at  
java.util.concurrent.ThreadPoolExecutor\$Worker.run(ThreadPoolExecutor.java:615) at java.lang.Thread.run(Thread.java:745)  
2015-03-13 20:17:24,867 [TezChild] INFO org.apache.hadoop.fs.s3native.NativeS3FileSystem - Received IOException while  
reading 'user/hadoop/tsato/readlarge/input/cloudian-s3.log.20141119', attempting to reopen. 2015-03-13 20:17:24,867 [TezChild]  
DEBUG org.jets3t.service.impl.rest.httpclient.HttpMethodReleaseInputStream - Released HttpMethod as its response data stream  
is fully consumed 2015-03-13 20:17:24,868 [TezChild] INFO org.apache.tez.dag.app.TaskAttemptListenerImpTezDag - Commit  
go/no-go request from attempt\_1426245338920\_0001\_1\_00\_000004\_0 2015-03-13 20:17:24,868 [TezChild] INFO  
org.apache.tez.dag.app.dag.impl.TaskImpl - attempt\_1426245338920\_0001\_1\_00\_000004\_0 given a go for committing the task  
output. {quote} It seems this is a regression, which was introduced by the following optimizations.  
<https://issues.apache.org/jira/browse/HADOOP-10589> <https://issues.apache.org/jira/browse/HADOOP-10457> Also, test cases  
should be reviewed so that it covers this scenario.

8. **summary:** Regression: s3n read failure recovery broken

**description:** s3n attempts to read again when it encounters IOException during read. But the current logic does not reopen the  
connection, thus, it ends up with no-op, and committing the wrong(truncated) output. Here's a stack trace as an example. {quote}  
2015-03-13 20:17:24,835 [TezChild] INFO org.apache.pig.backend.hadoop.executionengine.tez.runtime.PigProcessor - Starting  
output org.apache.tez.mapreduce.output.MROutput@52008dbd to vertex scope-12 2015-03-13 20:17:24,866 [TezChild] DEBUG  
org.jets3t.service.impl.rest.httpclient.HttpMethodReleaseInputStream - Released HttpMethod as its response data stream threw an  
exception org.apache.http.ConnectionClosedException: Premature end of Content-Length delimited message body (expected:  
296587138; received: 155648 at org.apache.http.impl.io.ContentLengthInputStream.read(ContentLengthInputStream.java:184) at  
org.apache.http.conn.EofSensorInputStream.read(EofSensorInputStream.java:138) at  
org.jets3t.service.io.InterruptableInputStream.read(InterruptableInputStream.java:78) at  
org.jets3t.service.impl.rest.httpclient.HttpMethodReleaseInputStream.read(HttpMethodReleaseInputStream.java:146) at  
org.apache.hadoop.fs.s3native.NativeS3FileSystem\$NativeS3FsInputStream.read(NativeS3FileSystem.java:145) at  
java.io.BufferedInputStream.read1(BufferedInputStream.java:273) at  
java.io.BufferedInputStream.read(BufferedInputStream.java:334) at java.io.DataInputStream.read(DataInputStream.java:100) at  
org.apache.hadoop.util.LineReader.fillBuffer(LineReader.java:180) at  
org.apache.hadoop.util.LineReader.readDefaultLine(LineReader.java:216) at  
org.apache.hadoop.util.LineReader.readLine(LineReader.java:174) at  
org.apache.hadoop.mapreduce.lib.input.LineRecordReader.nextKeyValue(LineRecordReader.java:185) at  
org.apache.pig.builtin.PigStorage.getNext(PigStorage.java:259) at  
org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.PigRecordReader.nextKeyValue(PigRecordReader.java:204) at  
org.apache.tez.mapreduce.lib.MRReaderMapReduce.next(MRReaderMapReduce.java:116) at  
org.apache.pig.backend.hadoop.executionengine.tez.plan.operator.POSimpleTezLoad.getNextTuple(POSimpleTezLoad.java:106)  
at org.apache.pig.backend.hadoop.executionengine.physicalLayer.PhysicalOperator.processInput(PhysicalOperator.java:307) at  
org.apache.pig.backend.hadoop.executionengine.physicalLayer.relationalOperators.POForEach.getNextTuple(POForEach.java:246)  
at org.apache.pig.backend.hadoop.executionengine.physicalLayer.PhysicalOperator.processInput(PhysicalOperator.java:307) at  
org.apache.pig.backend.hadoop.executionengine.physicalLayer.relationalOperators.POFilter.getNextTuple(POFilter.java:91) at  
org.apache.pig.backend.hadoop.executionengine.physicalLayer.PhysicalOperator.processInput(PhysicalOperator.java:307) at  
org.apache.pig.backend.hadoop.executionengine.tez.plan.operator.POStoreTez.getNextTuple(POStoreTez.java:117) at  
org.apache.pig.backend.hadoop.executionengine.tez.runtime.PigProcessor.runPipeline(PigProcessor.java:313) at  
org.apache.pig.backend.hadoop.executionengine.tez.runtime.PigProcessor.run(PigProcessor.java:192) at  
org.apache.tez.runtime.LogicalIOProcessorRuntimeTask.run(LogicalIOProcessorRuntimeTask.java:324) at  
org.apache.tez.runtime.task.TezTaskRunner\$TaskRunnerCallable\$1.run(TezTaskRunner.java:176) at  
org.apache.tez.runtime.task.TezTaskRunner\$TaskRunnerCallable\$1.run(TezTaskRunner.java:168) at  
java.security.AccessController.doPrivileged(Native Method) at javax.security.auth.Subject.doAs(Subject.java:415) at  
org.apache.hadoop.security.UserGroupInformation.doAs(UserGroupInformation.java:1628) at  
org.apache.tez.runtime.task.TezTaskRunner\$TaskRunnerCallable.call(TezTaskRunner.java:168) at  
org.apache.tez.runtime.task.TezTaskRunner\$TaskRunnerCallable.call(TezTaskRunner.java:163) at  
java.util.concurrent.FutureTask.run(FutureTask.java:262) at  
java.util.concurrent.ThreadPoolExecutor.runWorker(ThreadPoolExecutor.java:1145) at  
java.util.concurrent.ThreadPoolExecutor\$Worker.run(ThreadPoolExecutor.java:615) at java.lang.Thread.run(Thread.java:745)  
2015-03-13 20:17:24,867 [TezChild] INFO org.apache.hadoop.fs.s3native.NativeS3FileSystem - Received IOException while  
reading 'user/hadoop/tsato/readlarge/input/cloudian-s3.log.20141119', attempting to reopen. 2015-03-13 20:17:24,867 [TezChild]  
DEBUG org.jets3t.service.impl.rest.httpclient.HttpMethodReleaseInputStream - Released HttpMethod as its response data stream  
is fully consumed 2015-03-13 20:17:24,868 [TezChild] INFO org.apache.tez.dag.app.TaskAttemptListenerImpTezDag - Commit  
go/no-go request from attempt\_1426245338920\_0001\_1\_00\_000004\_0 2015-03-13 20:17:24,868 [TezChild] INFO  
org.apache.tez.dag.app.dag.impl.TaskImpl - attempt\_1426245338920\_0001\_1\_00\_000004\_0 given a go for committing the task  
output. {quote} It seems this is a regression, which was introduced by the following optimizations.  
<https://issues.apache.org/jira/browse/HADOOP-10589> <https://issues.apache.org/jira/browse/HADOOP-10457> Also, test cases  
should be reviewed so that it covers this scenario.

9. **summary:** Regression: s3n read failure recovery broken

**description:** s3n attempts to read again when it encounters IOException during read. But the current logic does not reopen the  
connection, thus, it ends up with no-op, and committing the wrong(truncated) output. Here's a stack trace as an example. {quote}  
2015-03-13 20:17:24,835 [TezChild] INFO org.apache.pig.backend.hadoop.executionengine.tez.runtime.PigProcessor - Starting  
output org.apache.tez.mapreduce.output.MROutput@52008dbd to vertex scope-12 2015-03-13 20:17:24,866 [TezChild] DEBUG  
org.jets3t.service.impl.rest.httpclient.HttpMethodReleaseInputStream - Released HttpMethod as its response data stream threw an  
exception org.apache.http.ConnectionClosedException: Premature end of Content-Length delimited message body (expected:  
296587138; received: 155648 at org.apache.http.impl.io.ContentLengthInputStream.read(ContentLengthInputStream.java:184) at

org.apache.http.conn.EofSensorInputStream.read(EofSensorInputStream.java:138) at  
org.jets3t.service.io.InterruptableInputStream.read(InterruptableInputStream.java:78) at  
org.jets3t.service.impl.rest.httpclient.HttpMethodReleaseInputStream.read(HttpMethodReleaseInputStream.java:146) at  
org.apache.hadoop.fs.s3native.NativeS3FileSystem\$NativeS3FsInputStream.read(NativeS3FileSystem.java:145) at  
java.io.BufferedInputStream.read1(BufferedInputStream.java:273) at  
java.io.BufferedInputStream.read(BufferedInputStream.java:334) at java.io.DataInputStream.read(DataInputStream.java:100) at  
org.apache.hadoop.util.LineReader.fillBuffer(LineReader.java:180) at  
org.apache.hadoop.util.LineReader.readDefaultLine(LineReader.java:216) at  
org.apache.hadoop.util.LineReader.readLine(LineReader.java:174) at  
org.apache.hadoop.mapreduce.lib.input.LineRecordReader.nextKeyValue(LineRecordReader.java:185) at  
org.apache.pig.builtin.PigStorage.getNext(PigStorage.java:259) at  
org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.PigRecordReader.nextKeyValue(PigRecordReader.java:204) at  
org.apache.tez.mapreduce.lib.MRReaderMapReduce.next(MRReaderMapReduce.java:116) at  
org.apache.pig.backend.hadoop.executionengine.tez.plan.operator.POSimpleTezLoad.getNextTuple(POSimpleTezLoad.java:106)  
at org.apache.pig.backend.hadoop.executionengine.physicalLayer.PhysicalOperator.processInput(PhysicalOperator.java:307) at  
org.apache.pig.backend.hadoop.executionengine.physicalLayer.relationalOperators.POForEach.getNextTuple(POForEach.java:246)  
at org.apache.pig.backend.hadoop.executionengine.physicalLayer.PhysicalOperator.processInput(PhysicalOperator.java:307) at  
org.apache.pig.backend.hadoop.executionengine.physicalLayer.relationalOperators.POFilter.getNextTuple(POFilter.java:91) at  
org.apache.pig.backend.hadoop.executionengine.physicalLayer.PhysicalOperator.processInput(PhysicalOperator.java:307) at  
org.apache.pig.backend.hadoop.executionengine.tez.plan.operator.POStoreTez.getNextTuple(POStoreTez.java:117) at  
org.apache.pig.backend.hadoop.executionengine.tez.runtime.PigProcessor.runPipeline(PigProcessor.java:313) at  
org.apache.pig.backend.hadoop.executionengine.tez.runtime.PigProcessor.run(PigProcessor.java:192) at  
org.apache.tez.runtime.LogicalIOProcessorRuntimeTask.run(LogicalIOProcessorRuntimeTask.java:324) at  
org.apache.tez.runtime.task.TezTaskRunner\$TaskRunnerCallable\$1.run(TezTaskRunner.java:176) at  
org.apache.tez.runtime.task.TezTaskRunner\$TaskRunnerCallable\$1.run(TezTaskRunner.java:168) at  
java.security.AccessController.doPrivileged(Native Method) at javax.security.auth.Subject.doAs(Subject.java:415) at  
org.apache.hadoop.security.UserGroupInformation.doAs(UserGroupInformation.java:1628) at  
org.apache.tez.runtime.task.TezTaskRunner\$TaskRunnerCallable.call(TezTaskRunner.java:168) at  
org.apache.tez.runtime.task.TezTaskRunner\$TaskRunnerCallable.call(TezTaskRunner.java:163) at  
java.util.concurrent.FutureTask.run(FutureTask.java:262) at  
java.util.concurrent.ThreadPoolExecutor.runWorker(ThreadPoolExecutor.java:1145) at  
java.util.concurrent.ThreadPoolExecutor\$Worker.run(ThreadPoolExecutor.java:615) at java.lang.Thread.run(Thread.java:745)  
2015-03-13 20:17:24,867 [TezChild] INFO org.apache.hadoop.fs.s3native.NativeS3FileSystem - Received IOException while  
reading 'user/hadoop/tsato/readlarge/input/cloudian-s3.log.20141119', attempting to reopen. 2015-03-13 20:17:24,867 [TezChild]  
DEBUG org.jets3t.service.impl.rest.httpclient.HttpMethodReleaseInputStream - Released HttpMethod as its response data stream  
is fully consumed 2015-03-13 20:17:24,868 [TezChild] INFO org.apache.tez.dag.app.TaskAttemptListenerImpTezDag - Commit  
go/no-go request from attempt\_1426245338920\_0001\_1\_00\_000004\_0 2015-03-13 20:17:24,868 [TezChild] INFO  
org.apache.tez.dag.app.dag.impl.TaskImpl - attempt\_1426245338920\_0001\_1\_00\_000004\_0 given a go for committing the task  
output. {quote} It seems this is a regression, which was introduced by the following optimizations.  
<https://issues.apache.org/jira/browse/HADOOP-10589> <https://issues.apache.org/jira/browse/HADOOP-10457> Also, test cases  
should be reviewed so that it covers this scenario.

## jira\_issues\_comments:

- body:** The first proposal without the test case. 2015-03-20 12:05:08,473 [TezChild] INFO  
org.apache.hadoop.fs.s3native.NativeS3FileSystem - Received IOException while reading  
'user/hadoop/tsato/readlarge/input/cloudian-s3.log.20141119', attempting to reopen. 2015-03-20 12:05:08,473 [TezChild] DEBUG  
org.jets3t.service.impl.rest.httpclient.RestStorageService - Retrieving All information for bucket shared and object  
user/hadoop/tsato/readlarge/input/cloudian-s3.log.20141119 Verified manually that it reopens a new connection after IOException.  
**label:** code-design
- The first patch with the updated test case.
- \\ \\ (x) \*{color:red}-1 overall{color}\* \\ \\ \\ Vote || Subsystem || Runtime || Comment || | {color:blue}0{color} | pre-patch | 14m  
32s | Pre-patch trunk compilation is healthy. || {color:green}+1{color} | @author | 0m 0s | The patch does not contain any  
@author tags. || {color:green}+1{color} | tests included | 0m 0s | The patch appears to include 1 new or modified test files. ||  
{color:red}-1{color} | whitespace | 0m 0s | The patch has 3 line(s) that end in whitespace. || {color:green}+1{color} | javac | 7m  
27s | There were no new javac warning messages. || {color:green}+1{color} | javadoc | 9m 36s | There were no new javadoc  
warning messages. || {color:green}+1{color} | release audit | 0m 23s | The applied patch does not increase the total number of  
release audit warnings. || {color:red}-1{color} | checkstyle | 5m 28s | The applied patch generated 1 additional checkstyle issues. |  
| {color:green}+1{color} | install | 1m 31s | mvn install still works. || {color:green}+1{color} | eclipse:eclipse | 0m 33s | The patch  
built with eclipse:eclipse. || {color:green}+1{color} | findbugs | 0m 38s | The patch does not introduce any new Findbugs (version  
2.0.3) warnings. || {color:green}+1{color} | tools/hadoop tests | 0m 15s | Tests passed in hadoop-aws. || | 40m 33s | \\ \\ \\  
Subsystem || Report/Notes || | Patch URL | <http://issues.apache.org/jira/secure/attachment/12705845/HADOOP-11730-branch-2.6.0.001.patch> || Optional Tests | javadoc javac unit findbugs checkstyle || git revision | trunk / 189a63a || whitespace |  
<https://builds.apache.org/job/PreCommit-HADOOP-Build/6167/artifact/patchprocess/whitespace.txt> || checkstyle |  
<https://builds.apache.org/job/PreCommit-HADOOP-Build/6167/artifact/patchprocess/checkstyle-result-diff.txt> || hadoop-aws test  
log | [https://builds.apache.org/job/PreCommit-HADOOP-Build/6167/artifact/patchprocess/testrun\\_hadoop-aws.txt](https://builds.apache.org/job/PreCommit-HADOOP-Build/6167/artifact/patchprocess/testrun_hadoop-aws.txt) || Test Results |  
<https://builds.apache.org/job/PreCommit-HADOOP-Build/6167/testReport/> || Console output |  
<https://builds.apache.org/job/PreCommit-HADOOP-Build/6167/console> | This message was automatically generated.
- +1 committing Here's the patch in sync with trunk; it also incorporates HADOOP-11851 in the close logic, as they go hand in  
hand. We can't have the recovery process damaged by ConnectionReset exceptions being picked up while it closes the old stream.
- body:** patch applied; tested against s3 EU. Given the nature of these problems, it may be good to start thinking about whether we  
can do things with better simulate failures; the test here is a good start, though we may want more complex policies...mockito

might be the tool to reach for.

**label:** test

6. FAILURE: Integrated in Hadoop-trunk-Commit #7653 (See [<https://builds.apache.org/job/Hadoop-trunk-Commit/7653/>]) HADOOP-11730. Regression: s3n read failure recovery broken. (Takenori Sato via stever) (stever: rev 19262d99ebbbd143a7ac9740d3a8e7c842b37591) \* hadoop-tools/hadoop-aws/src/main/java/org/apache/hadoop/fs/s3native/NativeS3FileSystem.java \* hadoop-tools/hadoop-aws/src/test/java/org/apache/hadoop/fs/s3native/NativeS3FileSystemContractBaseTest.java \* hadoop-common-project/hadoop-common/CHANGES.txt
7. \ \ \ (x) \* {color:red}-1 overall {color}\* | \ \ \ \ Vote | Subsystem | Runtime | Comment | | {color:blue}0 {color} | pre-patch | 14m 33s | Pre-patch trunk compilation is healthy. | | {color:green}+1 {color} | @author | 0m 0s | The patch does not contain any @author tags. | | {color:green}+1 {color} | tests included | 0m 0s | The patch appears to include 1 new or modified test files. | | {color:green}+1 {color} | whitespace | 0m 0s | The patch has no lines that end in whitespace. | | {color:green}+1 {color} | javac | 7m 27s | There were no new javac warning messages. | | {color:green}+1 {color} | javadoc | 9m 33s | There were no new javadoc warning messages. | | {color:green}+1 {color} | release audit | 0m 22s | The applied patch does not increase the total number of release audit warnings. | | {color:red}-1 {color} | checkstyle | 7m 43s | The applied patch generated 1 additional checkstyle issues. | | {color:green}+1 {color} | install | 1m 33s | mvn install still works. | | {color:green}+1 {color} | eclipse:eclipse | 0m 33s | The patch built with eclipse:eclipse. | | {color:green}+1 {color} | findbugs | 0m 38s | The patch does not introduce any new Findbugs (version 2.0.3) warnings. | | {color:green}+1 {color} | tools/hadoop tests | 0m 14s | Tests passed in hadoop-aws. | | | 42m 39s | \ \ \ \ Subsystem | Report/Notes | | Patch URL | <http://issues.apache.org/jira/secure/attachment/12727699/HADOOP-11730-002.patch> | | Optional Tests | javadoc javac unit findbugs checkstyle | | git revision | trunk / 416b843 | | checkstyle | <https://builds.apache.org/job/PreCommit-HADOOP-Build/6170/artifact/patchprocess/checkstyle-result-diff.txt> | | hadoop-aws test log | [https://builds.apache.org/job/PreCommit-HADOOP-Build/6170/artifact/patchprocess/testrun\\_hadoop-aws.txt](https://builds.apache.org/job/PreCommit-HADOOP-Build/6170/artifact/patchprocess/testrun_hadoop-aws.txt) | | Test Results | <https://builds.apache.org/job/PreCommit-HADOOP-Build/6170/testReport/> | | Console output | <https://builds.apache.org/job/PreCommit-HADOOP-Build/6170/console> | This message was automatically generated.
8. FAILURE: Integrated in Hadoop-Hdfs-trunk #2105 (See [<https://builds.apache.org/job/Hadoop-Hdfs-trunk/2105/>]) HADOOP-11730. Regression: s3n read failure recovery broken. (Takenori Sato via stever) (stever: rev 19262d99ebbbd143a7ac9740d3a8e7c842b37591) \* hadoop-common-project/hadoop-common/CHANGES.txt \* hadoop-tools/hadoop-aws/src/main/java/org/apache/hadoop/fs/s3native/NativeS3FileSystem.java \* hadoop-tools/hadoop-aws/src/test/java/org/apache/hadoop/fs/s3native/NativeS3FileSystemContractBaseTest.java
9. FAILURE: Integrated in Hadoop-Hdfs-trunk-Java8 #164 (See [<https://builds.apache.org/job/Hadoop-Hdfs-trunk-Java8/164/>]) HADOOP-11730. Regression: s3n read failure recovery broken. (Takenori Sato via stever) (stever: rev 19262d99ebbbd143a7ac9740d3a8e7c842b37591) \* hadoop-tools/hadoop-aws/src/test/java/org/apache/hadoop/fs/s3native/NativeS3FileSystemContractBaseTest.java \* hadoop-common-project/hadoop-common/CHANGES.txt \* hadoop-tools/hadoop-aws/src/main/java/org/apache/hadoop/fs/s3native/NativeS3FileSystem.java
10. FAILURE: Integrated in Hadoop-Yarn-trunk-Java8 #173 (See [<https://builds.apache.org/job/Hadoop-Yarn-trunk-Java8/173/>]) HADOOP-11730. Regression: s3n read failure recovery broken. (Takenori Sato via stever) (stever: rev 19262d99ebbbd143a7ac9740d3a8e7c842b37591) \* hadoop-common-project/hadoop-common/CHANGES.txt \* hadoop-tools/hadoop-aws/src/test/java/org/apache/hadoop/fs/s3native/NativeS3FileSystemContractBaseTest.java \* hadoop-tools/hadoop-aws/src/main/java/org/apache/hadoop/fs/s3native/NativeS3FileSystem.java
11. SUCCESS: Integrated in Hadoop-Yarn-trunk #907 (See [<https://builds.apache.org/job/Hadoop-Yarn-trunk/907/>]) HADOOP-11730. Regression: s3n read failure recovery broken. (Takenori Sato via stever) (stever: rev 19262d99ebbbd143a7ac9740d3a8e7c842b37591) \* hadoop-tools/hadoop-aws/src/main/java/org/apache/hadoop/fs/s3native/NativeS3FileSystem.java \* hadoop-common-project/hadoop-common/CHANGES.txt \* hadoop-tools/hadoop-aws/src/test/java/org/apache/hadoop/fs/s3native/NativeS3FileSystemContractBaseTest.java
12. FAILURE: Integrated in Hadoop-Mapreduce-trunk-Java8 #174 (See [<https://builds.apache.org/job/Hadoop-Mapreduce-trunk-Java8/174/>]) HADOOP-11730. Regression: s3n read failure recovery broken. (Takenori Sato via stever) (stever: rev 19262d99ebbbd143a7ac9740d3a8e7c842b37591) \* hadoop-tools/hadoop-aws/src/test/java/org/apache/hadoop/fs/s3native/NativeS3FileSystemContractBaseTest.java \* hadoop-tools/hadoop-aws/src/main/java/org/apache/hadoop/fs/s3native/NativeS3FileSystem.java \* hadoop-common-project/hadoop-common/CHANGES.txt
13. FAILURE: Integrated in Hadoop-Mapreduce-trunk #2123 (See [<https://builds.apache.org/job/Hadoop-Mapreduce-trunk/2123/>]) HADOOP-11730. Regression: s3n read failure recovery broken. (Takenori Sato via stever) (stever: rev 19262d99ebbbd143a7ac9740d3a8e7c842b37591) \* hadoop-tools/hadoop-aws/src/test/java/org/apache/hadoop/fs/s3native/NativeS3FileSystemContractBaseTest.java \* hadoop-common-project/hadoop-common/CHANGES.txt \* hadoop-tools/hadoop-aws/src/main/java/org/apache/hadoop/fs/s3native/NativeS3FileSystem.java
14. Pulled this into 2.6.1 after [~ajisakaa] verified that the patch applies cleanly. Ran compilation and TestInMemoryNativeS3FileSystemContract before the push.