

git_comments:

1. Check with the user before blowing away data.
2. Force the active to roll its log
3. Download that checkpoint into our storage directories.
4. * * Tool which allows the standby node's storage directories to be bootstrapped * by copying the latest namespace snapshot from the active namenode. This is * used when first configuring an HA cluster.
5. Load the newly formatted image, using all of the directories (including shared edits)
6. Format the storage (writes VERSION file)
7. Ensure that we have enough edits already in the shared directory to start up from the last checkpoint on the active.
8. * * Licensed to the Apache Software Foundation (ASF) under one * or more contributor license agreements. See the NOTICE file * distributed with this work for additional information * regarding copyright ownership. The ASF licenses this file * to you under the Apache License, Version 2.0 (the * "License"); you may not use this file except in compliance * with the License. You may obtain a copy of the License at * * <http://www.apache.org/licenses/LICENSE-2.0> * * Unless required by applicable law or agreed to in writing, software * distributed under the License is distributed on an "AS IS" BASIS, * WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied. * See the License for the specific language governing permissions and * limitations under the License.
9. Delete the segment.
10. Roll edit logs a few times to inflate txid
11. * * Test for the case where the shared edits dir doesn't have * all of the recent edit logs.
12. Should have copied over the namespace from the active
13. Should pass with -force
14. Should not pass since standby dirs exist, force not given
15. Trying to bootstrap standby should now fail since the edit logs aren't available in the shared dir.
16. * * Test for the base success case. The primary NN * hasn't made any checkpoints, and we copy the fsimage_0 * file over and start up.
17. Make checkpoint
18. Should have created edits_1-2 in shared edits dir
19. We should now be able to start the standby successfully.
20. * * Licensed to the Apache Software Foundation (ASF) under one * or more contributor license agreements. See the NOTICE file * distributed with this work for additional information * regarding copyright ownership. The ASF licenses this file * to you under the Apache License, Version 2.0 (the * "License"); you may not use this file except in compliance * with the License. You may obtain a copy of the License at * * <http://www.apache.org/licenses/LICENSE-2.0> * * Unless required by applicable law or agreed to in writing, software * distributed under the License is distributed on an "AS IS" BASIS, * WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied. * See the License for the specific language governing permissions and * limitations under the License.
21. * * Test for downloading a checkpoint made at a later checkpoint * from the active.
22. * * Check whether the given storage directories already exist. * If running in interactive mode, will prompt the user for each * directory to allow them to format anyway. Otherwise, returns * false, unless 'force' is specified. * * @param dirsToFormat the dirs to check * @param force format regardless of whether dirs exist * @param interactive prompt the user when a dir exists * @return true if formatting should proceed * @throws IOException
23. Don't confirm, always format.
24. Don't ask - always don't format
25. avoid warning
26. aborted

git_commits:

1. **summary:** HDFS-2731. Add command to bootstrap the Standby Node's name directories from the Active NameNode. Contributed by Todd Lipcon.
message: HDFS-2731. Add command to bootstrap the Standby Node's name directories from the Active NameNode. Contributed by Todd Lipcon. git-svn-id:
<https://svn.apache.org/repos/asf/hadoop/common/branches/branch-0.23@1299809> 13f79535-47bb-0310-9956-ffa450edef68

github_issues:

github_issues_comments:

github_pulls:

github_pulls_comments:

github_pulls_reviews:

jira_issues:

1. **summary:** HA: Autopopulate standby name dirs if they're empty
description: To setup a SBN we currently format the primary then manually copy the name dirs to the SBN. The SBN should do this automatically. Specifically, on NN startup, if HA with a shared edits dir is configured and populated, if the SBN has empty name dirs it should download the image and log from the primary (as an optimization it could copy the logs from the shared dir). If the other NN is still in standby then it should fail to start as it does currently.
label: code-design
2. **summary:** HA: Autopopulate standby name dirs if they're empty
description: To setup a SBN we currently format the primary then manually copy the name dirs to the SBN. The SBN should do this automatically. Specifically, on NN startup, if HA with a shared edits dir is configured and populated, if the SBN has empty name dirs it should download the image and log from the primary (as an optimization it could copy the logs from the shared dir). If the other NN is still in standby then it should fail to start as it does currently.
3. **summary:** HA: Autopopulate standby name dirs if they're empty
description: To setup a SBN we currently format the primary then manually copy the name dirs to the SBN. The SBN should do this automatically. Specifically, on NN startup, if HA with a shared edits dir is configured and populated, if the SBN has empty name dirs it should download the image and log from the primary (as an optimization it could copy the logs from the shared dir). If the other NN is still in standby then it should fail to start as it does currently.
4. **summary:** HA: Autopopulate standby name dirs if they're empty
description: To setup a SBN we currently format the primary then manually copy the name dirs to the SBN. The SBN should do this automatically. Specifically, on NN startup, if HA with a shared edits dir is configured and populated, if the SBN has empty name dirs it should download the image and log from the primary (as an optimization it could copy the logs from the shared dir). If the other NN is still in standby then it should fail to start as it does currently.
5. **summary:** HA: Autopopulate standby name dirs if they're empty
description: To setup a SBN we currently format the primary then manually copy the name dirs to the SBN. The SBN should do this automatically. Specifically, on NN startup, if HA with a shared edits dir is configured and populated, if the SBN has empty name dirs it should download the image and log from the primary (as an optimization it could copy the logs from the shared dir). If the other NN is still in standby then it should fail to start as it does currently.
label: code-design
6. **summary:** HA: Autopopulate standby name dirs if they're empty
description: To setup a SBN we currently format the primary then manually copy the name dirs to the SBN. The SBN should do this automatically. Specifically, on NN startup, if HA with a shared edits dir is configured and populated, if the SBN has empty name dirs it should download the image and log from the primary (as an optimization it could copy the logs from the shared dir). If the other NN is still in standby then it should fail to start as it does currently.
7. **summary:** HA: Autopopulate standby name dirs if they're empty
description: To setup a SBN we currently format the primary then manually copy the name dirs to the SBN. The SBN should do this automatically. Specifically, on NN startup, if HA with a shared edits dir is configured and populated, if the SBN has empty name dirs it should download the image and log from the primary (as an optimization it could copy the logs from the shared dir). If the other NN is still in standby then it should fail to start as it does currently.
8. **summary:** HA: Autopopulate standby name dirs if they're empty
description: To setup a SBN we currently format the primary then manually copy the name dirs to the SBN. The SBN should do this automatically. Specifically, on NN startup, if HA with a shared edits dir is

configured and populated, if the SBN has empty name dirs it should download the image and log from the primary (as an optimization it could copy the logs from the shared dir). If the other NN is still in standby then it should fail to start as it does currently.

label: code-design

- [illegible]

description: To setup a SBN we currently format the primary then manually copy the name dirs to the SBN. The SBN should do this automatically. Specifically, on NN startup, if HA with a shared edits dir is configured and populated, if the SBN has empty name dirs it should download the image and log from the primary (as an optimization it could copy the logs from the shared dir). If the other NN is still in standby then it should fail to start as it does currently.

19. **summary:** HA: Autopopulate standby name dirs if they're empty

description: To setup a SBN we currently format the primary then manually copy the name dirs to the SBN. The SBN should do this automatically. Specifically, on NN startup, if HA with a shared edits dir is configured and populated, if the SBN has empty name dirs it should download the image and log from the primary (as an optimization it could copy the logs from the shared dir). If the other NN is still in standby then it should fail to start as it does currently.

label: code-design

20. **summary:** HA: Autopopulate standby name dirs if they're empty

description: To setup a SBN we currently format the primary then manually copy the name dirs to the SBN. The SBN should do this automatically. Specifically, on NN startup, if HA with a shared edits dir is configured and populated, if the SBN has empty name dirs it should download the image and log from the primary (as an optimization it could copy the logs from the shared dir). If the other NN is still in standby then it should fail to start as it does currently.

label: code-design

21. **summary:** HA: Autopopulate standby name dirs if they're empty

description: To setup a SBN we currently format the primary then manually copy the name dirs to the SBN. The SBN should do this automatically. Specifically, on NN startup, if HA with a shared edits dir is configured and populated, if the SBN has empty name dirs it should download the image and log from the primary (as an optimization it could copy the logs from the shared dir). If the other NN is still in standby then it should fail to start as it does currently.

22. **summary:** HA: Autopopulate standby name dirs if they're empty

description: To setup a SBN we currently format the primary then manually copy the name dirs to the SBN. The SBN should do this automatically. Specifically, on NN startup, if HA with a shared edits dir is configured and populated, if the SBN has empty name dirs it should download the image and log from the primary (as an optimization it could copy the logs from the shared dir). If the other NN is still in standby then it should fail to start as it does currently.

23. **summary:** HA: Autopopulate standby name dirs if they're empty

description: To setup a SBN we currently format the primary then manually copy the name dirs to the SBN. The SBN should do this automatically. Specifically, on NN startup, if HA with a shared edits dir is configured and populated, if the SBN has empty name dirs it should download the image and log from the primary (as an optimization it could copy the logs from the shared dir). If the other NN is still in standby then it should fail to start as it does currently.

24. **summary:** HA: Autopopulate standby name dirs if they're empty

description: To setup a SBN we currently format the primary then manually copy the name dirs to the SBN. The SBN should do this automatically. Specifically, on NN startup, if HA with a shared edits dir is configured and populated, if the SBN has empty name dirs it should download the image and log from the primary (as an optimization it could copy the logs from the shared dir). If the other NN is still in standby then it should fail to start as it does currently.

25. **summary:** HA: Autopopulate standby name dirs if they're empty

description: To setup a SBN we currently format the primary then manually copy the name dirs to the SBN. The SBN should do this automatically. Specifically, on NN startup, if HA with a shared edits dir is configured and populated, if the SBN has empty name dirs it should download the image and log from the primary (as an optimization it could copy the logs from the shared dir). If the other NN is still in standby then it should fail to start as it does currently.

jira_issues_comments:

1. **body:** bq. as an optimization it could copy the logs from the shared dir I don't think it's necessarily an optimization - might actually be _easier_ to implement this way :) bq. If the other NN is still in standby then it should fail to start as it does currently Can you explain what you mean by this? Why not allow it to download the image from the other NN anyway?
label: code-design
2. Wrt #1, if we get the image from the NN and the edits from the shared dir, are we sure they'll always match, eg what if we're rolling at the same time (the other NN could be primary and active)? I was

- thinking asking for both from the primary would mean we always get matched sets and therefore don't need to worry about races. Wrt #2, yea, I thinking we should be explicit (don't have to worry about eg the shared dir being populated by neither NN having populated name dirs, which we know won't be the case if the other is active), but on 2nd thought I think your suggestion is better.
3. The primary shouldn't be removing any old images unless it's taking checkpoints. But there won't be checkpoints if the standby isn't running yet (assuming the standby is the one doing checkpointing). So if we get the most recent image from the NN, then we should always have enough edits in the shared dir to roll forward from there.
 4. Don't we support using both the SBN and 2NN though? There are issues requiring the SBN be the checkpointer, eg no one is doing checkpointing if the primary dies and you fail-over or the SBN still checkpoints but doesn't have enough memory.
 5. **body:** We don't currently support using the 2NN with HA -- we'd need to improve it to upload the checkpointed images to both NNs. If we think this is critical we should file another JIRA for it... in HDFS-2291 I actually have added a check to the 2NN startup that makes it fail to start if HA is enabled.
label: code-design
 6. Thanks. Filed HDFS-2736 to figure it out.
 7. >if the SBN has empty name dirs it should download the image and log from the primary (as an optimization it could copy the logs from the shared dir). I am missing this: both Image and Edits should be on shared dirs (ie no need to copy image from primary - it should be available in shared dir).
 8. **body:** Can we generalize this: on startup, a NN (PNN or SBN) should make all dirs consistent - i.e. copy any missing images and logs.
label: code-design
 9. bq. I am missing this: both Image and Edits should be on shared dirs (ie no need to copy image from primary - it should be available in shared dir). Since we generalized the edit log interface, we can't store images on the shared edit storage. We only get the stream abstraction for it. We could also extend the system to require a shared *image* dir, but then we couldn't support BK-based HA. (In the future when we can store checkpoints in HDFS itself we wouldn't have that issue)
 10. I dislike the idea of pulling the image from the primary through a protocol. (But I guess it is needed for the BNN solution.) For shared storage one could require that a shared dir be configured (NFS now and HDFS in the future by the time we get Bookkeeper.).
 11. Converting to top-level issue with commit of HDFS-1623.
 12. Attached patch adds a new command `{ {hadoop namenode -bootstrapStandby} }`. This command connects to the running active, fetches nameservice info, displays it for the user, and allows the user to reformat the standby's local directories based on the remote NN's info. Then, it uses the existing checkpoint transfer code to copy the most recent image from the standby into the local directories. It also ensures that the remote node's shared edits dir is visible on the local machine before performing the reformat.
 13. Here's an example output: `{code}`

```
===== About to bootstrap Standby ID
nn2 from: Nameservice ID: nameserviceId1 Other Namenode ID: nn1 Other NN's HTTP address: todd-
w510:50081 Other NN's IPC address: todd-w510/127.0.0.1:8021 Namespace ID: 1273108605 Block pool
ID: BP-2075569167-127.0.0.1-1330663690093 Cluster ID: CID-d2c547b6-1a29-4bd9-a1fa-
d5fdb8120f1 Layout version: -40
===== Re-format filesystem in
/tmp/name2-name ? (Y or N) y Re-format filesystem in /tmp/name2-name2 ? (Y or N) y 12/03/05
14:34:53 INFO namenode.NNStorage: Storage directory /tmp/name2-name has been successfully
formatted. 12/03/05 14:34:53 INFO namenode.NNStorage: Storage directory /tmp/name2-name2 has
been successfully formatted. 12/03/05 14:34:54 INFO namenode.TransferFsImage: Downloaded file
fsimage.ckpt_0000000000000000026 size 143 bytes. 12/03/05 14:34:54 INFO namenode.NameNode:
SHUTDOWN_MSG: /*****
SHUTDOWN_MSG: Shutting down NameNode at todd-w510/127.0.0.1
*****/ {code}
```
 14. -1 overall. Here are the results of testing the latest attachment
<http://issues.apache.org/jira/secure/attachment/12517142/hdfs-2731.txt> against trunk revision . +1
@author. The patch does not contain any @author tags. +1 tests included. The patch appears to include 8
new or modified tests. -1 patch. The patch command could not apply the patch. Console output:
<https://builds.apache.org/job/PreCommit-HDFS-Build/1955//console> This message is automatically
generated.
 15. Patch failed to apply since I accidentally changed SecurityUtil in this HDFS patch. See HADOOP-8141.

16. I've just kicked Jenkins to run test-patch for this again now that HADOOP-8141 is committed.
17. -1 overall. Here are the results of testing the latest attachment
<http://issues.apache.org/jira/secure/attachment/12517142/hdfs-2731.txt> against trunk revision . +1
@author. The patch does not contain any @author tags. +1 tests included. The patch appears to include 8 new or modified tests. -1 patch. The patch command could not apply the patch. Console output:
<https://builds.apache.org/job/PreCommit-HDFS-Build/1965//console> This message is automatically generated.
18. Of course test-patch is going to fail since the patches conflict. :) Here's a patch I'm uploading on behalf of Todd which removes the changes from hadoop-common-project.
19. **body:** Patch looks pretty good, Todd. Just two tiny nits. +1 once these are addressed and Jenkins comes back clean. # Put a space before the "+" - System.err.println("Format aborted in "+ curDir); # Add a class comment for BootstrapStandby.
label: code-design
20. **body:** Updated patch: - fixes style nit - adds javadoc and interface audience annotation (Private) to the new public class
label: code-design
21. -1 overall. Here are the results of testing the latest attachment
<http://issues.apache.org/jira/secure/attachment/12517466/hdfs-2731.txt> against trunk revision . +1
@author. The patch does not contain any @author tags. +1 tests included. The patch appears to include 8 new or modified tests. +1 javadoc. The javadoc tool did not generate any warning messages. +1 javac. The applied patch does not increase the total number of javac compiler warnings. +1 eclipse:eclipse. The patch built with eclipse:eclipse. +1 findbugs. The patch does not introduce any new Findbugs (version 1.3.9) warnings. +1 release audit. The applied patch does not increase the total number of release audit warnings. -1 core tests. The patch failed these unit tests:
org.apache.hadoop.hdfs.server.common.TestDistributedUpgrade +1 contrib tests. The patch passed contrib unit tests. Test results: <https://builds.apache.org/job/PreCommit-HDFS-Build/1966//testReport/>
Console output: <https://builds.apache.org/job/PreCommit-HDFS-Build/1966//console> This message is automatically generated.
22. -1 overall. Here are the results of testing the latest attachment
<http://issues.apache.org/jira/secure/attachment/12517508/hdfs-2731.txt> against trunk revision . +1
@author. The patch does not contain any @author tags. +1 tests included. The patch appears to include 8 new or modified tests. +1 javadoc. The javadoc tool did not generate any warning messages. +1 javac. The applied patch does not increase the total number of javac compiler warnings. +1 eclipse:eclipse. The patch built with eclipse:eclipse. +1 findbugs. The patch does not introduce any new Findbugs (version 1.3.9) warnings. +1 release audit. The applied patch does not increase the total number of release audit warnings. -1 core tests. The patch failed these unit tests:
org.apache.hadoop.hdfs.server.blockmanagement.TestBlocksWithNotEnoughRacks +1 contrib tests. The patch passed contrib unit tests. Test results: <https://builds.apache.org/job/PreCommit-HDFS-Build/1969//testReport/> Console output: <https://builds.apache.org/job/PreCommit-HDFS-Build/1969//console> This message is automatically generated.
23. Committed to 23.3 and trunk, thanks for the reviews Aaron.
24. Integrated in Hadoop-Hdfs-trunk-Commit #1943 (See [<https://builds.apache.org/job/Hadoop-Hdfs-trunk-Commit/1943/>]) HDFS-2731. Add command to bootstrap the Standby Node's name directories from the Active NameNode. Contributed by Todd Lipcon. (Revision 1299807) Result = SUCCESS todd :
<http://svn.apache.org/viewcvcs.cgi/?root=Apache-SVN&view=rev&rev=1299807> Files : *
/hadoop/common/trunk/hadoop-hdfs-project/hadoop-hdfs/CHANGES.txt *
/hadoop/common/trunk/hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/common/HdfsServerConstants.java *
/hadoop/common/trunk/hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/FSImage.java *
/hadoop/common/trunk/hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/FSNamesystem.java *
/hadoop/common/trunk/hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/NNStorage.java *
/hadoop/common/trunk/hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/NameNode.java *
/hadoop/common/trunk/hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/NameNodeRpcServer.java *
/hadoop/common/trunk/hadoop-hdfs-project/hadoop-

- hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/SecondaryNameNode.java *
- /hadoop/common/trunk/hadoop-hdfs-project/hadoop-
- hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/TransferFsImage.java *
- /hadoop/common/trunk/hadoop-hdfs-project/hadoop-
- hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/ha/BootstrapStandby.java *
- /hadoop/common/trunk/hadoop-hdfs-project/hadoop-
- hdfs/src/main/java/org/apache/hadoop/hdfs/tools/DFSck.java * /hadoop/common/trunk/hadoop-hdfs-
- project/hadoop-hdfs/src/test/java/org/apache/hadoop/hdfs/server/namenode/TestEditLog.java *
- /hadoop/common/trunk/hadoop-hdfs-project/hadoop-
- hdfs/src/test/java/org/apache/hadoop/hdfs/server/namenode/ha/TestBootstrapStandby.java *
- /hadoop/common/trunk/hadoop-hdfs-project/hadoop-
- hdfs/src/test/java/org/apache/hadoop/test/GenericTestUtils.java
25. Integrated in Hadoop-Common-trunk-Commit #1868 (See [<https://builds.apache.org/job/Hadoop-Common-trunk-Commit/1868/>]) HDFS-2731. Add command to bootstrap the Standby Node's name directories from the Active NameNode. Contributed by Todd Lipcon. (Revision 1299807) Result = SUCCESS todd : <http://svn.apache.org/viewcvcs.cgi/?root=Apache-SVN&view=rev&rev=1299807> Files :
- * /hadoop/common/trunk/hadoop-hdfs-project/hadoop-hdfs/CHANGES.txt *
- /hadoop/common/trunk/hadoop-hdfs-project/hadoop-
- hdfs/src/main/java/org/apache/hadoop/hdfs/server/common/HdfsServerConstants.java *
- /hadoop/common/trunk/hadoop-hdfs-project/hadoop-
- hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/FsImage.java *
- /hadoop/common/trunk/hadoop-hdfs-project/hadoop-
- hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/FsNamesystem.java *
- /hadoop/common/trunk/hadoop-hdfs-project/hadoop-
- hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/NNStorage.java *
- /hadoop/common/trunk/hadoop-hdfs-project/hadoop-
- hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/NameNode.java *
- /hadoop/common/trunk/hadoop-hdfs-project/hadoop-
- hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/NameNodeRpcServer.java *
- /hadoop/common/trunk/hadoop-hdfs-project/hadoop-
- hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/SecondaryNameNode.java *
- /hadoop/common/trunk/hadoop-hdfs-project/hadoop-
- hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/TransferFsImage.java *
- /hadoop/common/trunk/hadoop-hdfs-project/hadoop-
- hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/ha/BootstrapStandby.java *
- /hadoop/common/trunk/hadoop-hdfs-project/hadoop-
- hdfs/src/main/java/org/apache/hadoop/hdfs/tools/DFSck.java * /hadoop/common/trunk/hadoop-hdfs-
- project/hadoop-hdfs/src/test/java/org/apache/hadoop/hdfs/server/namenode/TestEditLog.java *
- /hadoop/common/trunk/hadoop-hdfs-project/hadoop-
- hdfs/src/test/java/org/apache/hadoop/hdfs/server/namenode/ha/TestBootstrapStandby.java *
- /hadoop/common/trunk/hadoop-hdfs-project/hadoop-
- hdfs/src/test/java/org/apache/hadoop/test/GenericTestUtils.java
26. Integrated in Hadoop-Hdfs-0.23-Commit #665 (See [<https://builds.apache.org/job/Hadoop-Hdfs-0.23-Commit/665/>]) HDFS-2731. Add command to bootstrap the Standby Node's name directories from the Active NameNode. Contributed by Todd Lipcon. (Revision 1299809) Result = SUCCESS todd :
- <http://svn.apache.org/viewcvcs.cgi/?root=Apache-SVN&view=rev&rev=1299809> Files : *
- /hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-hdfs/CHANGES.txt *
- /hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-
- hdfs/src/main/java/org/apache/hadoop/hdfs/server/common/HdfsServerConstants.java *
- /hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-
- hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/FsImage.java *
- /hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-
- hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/FsNamesystem.java *
- /hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-
- hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/NNStorage.java *
- /hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-
- hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/NameNode.java *
- /hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-
- hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/NameNodeRpcServer.java *

/hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/SecondaryNameNode.java *
 /hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/TransferFsImage.java *
 /hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/ha/BootstrapStandby.java *
 /hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/tools/DFSck.java * /hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-hdfs/src/test/java/org/apache/hadoop/hdfs/server/namenode/TestEditLog.java *
 /hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-hdfs/src/test/java/org/apache/hadoop/hdfs/server/namenode/ha/TestBootstrapStandby.java *
 /hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-hdfs/src/test/java/org/apache/hadoop/test/GenericTestUtils.java

27. Integrated in Hadoop-Common-0.23-Commit #674 (See [<https://builds.apache.org/job/Hadoop-Common-0.23-Commit/674/>]) HDFS-2731. Add command to bootstrap the Standby Node's name directories from the Active NameNode. Contributed by Todd Lipcon. (Revision 1299809) Result = SUCCESS todd : <http://svn.apache.org/viewcvcs.cgi/?root=Apache-SVN&view=rev&rev=1299809> Files : *
 /hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-hdfs/CHANGES.txt *
 /hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/common/HdfsServerConstants.java *
 /hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/FsImage.java *
 /hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/FsNamesystem.java *
 /hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/NNStorage.java *
 /hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/NameNode.java *
 /hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/NameNodeRpcServer.java *
 /hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/SecondaryNameNode.java *
 /hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/TransferFsImage.java *
 /hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/ha/BootstrapStandby.java *
 /hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/tools/DFSck.java * /hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-hdfs/src/test/java/org/apache/hadoop/hdfs/server/namenode/TestEditLog.java *
 /hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-hdfs/src/test/java/org/apache/hadoop/hdfs/server/namenode/ha/TestBootstrapStandby.java *
 /hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-hdfs/src/test/java/org/apache/hadoop/test/GenericTestUtils.java

28. Integrated in Hadoop-Mapreduce-trunk-Commit #1877 (See [<https://builds.apache.org/job/Hadoop-Mapreduce-trunk-Commit/1877/>]) HDFS-2731. Add command to bootstrap the Standby Node's name directories from the Active NameNode. Contributed by Todd Lipcon. (Revision 1299807) Result = ABORTED todd : <http://svn.apache.org/viewcvcs.cgi/?root=Apache-SVN&view=rev&rev=1299807> Files : *
 /hadoop/common/trunk/hadoop-hdfs-project/hadoop-hdfs/CHANGES.txt *
 /hadoop/common/trunk/hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/common/HdfsServerConstants.java *
 /hadoop/common/trunk/hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/FsImage.java *
 /hadoop/common/trunk/hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/FsNamesystem.java *
 /hadoop/common/trunk/hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/NNStorage.java *
 /hadoop/common/trunk/hadoop-hdfs-project/hadoop-

- hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/NameNode.java *
- /hadoop/common/trunk/hadoop-hdfs-project/hadoop-
- hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/NameNodeRpcServer.java *
- /hadoop/common/trunk/hadoop-hdfs-project/hadoop-
- hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/SecondaryNameNode.java *
- /hadoop/common/trunk/hadoop-hdfs-project/hadoop-
- hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/TransferFsImage.java *
- /hadoop/common/trunk/hadoop-hdfs-project/hadoop-
- hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/ha/BootstrapStandby.java *
- /hadoop/common/trunk/hadoop-hdfs-project/hadoop-
- hdfs/src/main/java/org/apache/hadoop/hdfs/tools/DFSck.java * /hadoop/common/trunk/hadoop-hdfs-
- project/hadoop-hdfs/src/test/java/org/apache/hadoop/hdfs/server/namenode/TestEditLog.java *
- /hadoop/common/trunk/hadoop-hdfs-project/hadoop-
- hdfs/src/test/java/org/apache/hadoop/hdfs/server/namenode/ha/TestBootstrapStandby.java *
- /hadoop/common/trunk/hadoop-hdfs-project/hadoop-
- hdfs/src/test/java/org/apache/hadoop/test/GenericTestUtils.java
29. Integrated in Hadoop-Mapreduce-0.23-Commit #682 (See [<https://builds.apache.org/job/Hadoop-Mapreduce-0.23-Commit/682/>]) HDFS-2731. Add command to bootstrap the Standby Node's name directories from the Active NameNode. Contributed by Todd Lipcon. (Revision 1299809) Result = ABORTED todd : <http://svn.apache.org/viewcvcs.cgi/?root=Apache-SVN&view=rev&rev=1299809> Files : * /hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-hdfs/CHANGES.txt *
- /hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-
- hdfs/src/main/java/org/apache/hadoop/hdfs/server/common/HdfsServerConstants.java *
- /hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-
- hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/FsImage.java *
- /hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-
- hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/FsNamesystem.java *
- /hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-
- hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/NNStorage.java *
- /hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-
- hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/NameNode.java *
- /hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-
- hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/NameNodeRpcServer.java *
- /hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-
- hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/SecondaryNameNode.java *
- /hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-
- hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/TransferFsImage.java *
- /hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-
- hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/ha/BootstrapStandby.java *
- /hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-
- hdfs/src/main/java/org/apache/hadoop/hdfs/tools/DFSck.java * /hadoop/common/branches/branch-
- 0.23/hadoop-hdfs-project/hadoop-
- hdfs/src/test/java/org/apache/hadoop/hdfs/server/namenode/TestEditLog.java *
- /hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-
- hdfs/src/test/java/org/apache/hadoop/hdfs/server/namenode/ha/TestBootstrapStandby.java *
- /hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-
- hdfs/src/test/java/org/apache/hadoop/test/GenericTestUtils.java
30. Integrated in Hadoop-Mapreduce-0.23-Build #224 (See [<https://builds.apache.org/job/Hadoop-Mapreduce-0.23-Build/224/>]) HDFS-2731. Add command to bootstrap the Standby Node's name directories from the Active NameNode. Contributed by Todd Lipcon. (Revision 1299809) Result = FAILURE todd : <http://svn.apache.org/viewcvcs.cgi/?root=Apache-SVN&view=rev&rev=1299809> Files :
- * /hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-hdfs/CHANGES.txt *
- /hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-
- hdfs/src/main/java/org/apache/hadoop/hdfs/server/common/HdfsServerConstants.java *
- /hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-
- hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/FsImage.java *
- /hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-
- hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/FsNamesystem.java *
- /hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-

- hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/NNStorage.java *
- /hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/NameNode.java *
- /hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/NameNodeRpcServer.java *
- /hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/SecondaryNameNode.java *
- /hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/TransferFsImage.java *
- /hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/ha/BootstrapStandby.java *
- /hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/tools/DFSck.java * /hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-hdfs/src/test/java/org/apache/hadoop/hdfs/server/namenode/TestEditLog.java *
- /hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-hdfs/src/test/java/org/apache/hadoop/hdfs/server/namenode/ha/TestBootstrapStandby.java *
- /hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-hdfs/src/test/java/org/apache/hadoop/test/GenericTestUtils.java
31. Integrated in Hadoop-Mapreduce-trunk #1018 (See [<https://builds.apache.org/job/Hadoop-Mapreduce-trunk/1018/>]) HDFS-2731. Add command to bootstrap the Standby Node's name directories from the Active NameNode. Contributed by Todd Lipcon. (Revision 1299807) Result = SUCCESS todd : <http://svn.apache.org/viewcvcs.cgi/?root=Apache-SVN&view=rev&rev=1299807> Files : *
- /hadoop/common/trunk/hadoop-hdfs-project/hadoop-hdfs/CHANGES.txt *
- /hadoop/common/trunk/hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/common/HdfsServerConstants.java *
- /hadoop/common/trunk/hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/FsImage.java *
- /hadoop/common/trunk/hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/FsNamesystem.java *
- /hadoop/common/trunk/hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/NNStorage.java *
- /hadoop/common/trunk/hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/NameNode.java *
- /hadoop/common/trunk/hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/NameNodeRpcServer.java *
- /hadoop/common/trunk/hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/SecondaryNameNode.java *
- /hadoop/common/trunk/hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/TransferFsImage.java *
- /hadoop/common/trunk/hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/ha/BootstrapStandby.java *
- /hadoop/common/trunk/hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/tools/DFSck.java * /hadoop/common/trunk/hadoop-hdfs-project/hadoop-hdfs/src/test/java/org/apache/hadoop/hdfs/server/namenode/TestEditLog.java *
- /hadoop/common/trunk/hadoop-hdfs-project/hadoop-hdfs/src/test/java/org/apache/hadoop/hdfs/server/namenode/ha/TestBootstrapStandby.java *
- /hadoop/common/trunk/hadoop-hdfs-project/hadoop-hdfs/src/test/java/org/apache/hadoop/test/GenericTestUtils.java
32. Integrated in Hadoop-Hdfs-trunk #983 (See [<https://builds.apache.org/job/Hadoop-Hdfs-trunk/983/>]) HDFS-2731. Add command to bootstrap the Standby Node's name directories from the Active NameNode. Contributed by Todd Lipcon. (Revision 1299807) Result = FAILURE todd : <http://svn.apache.org/viewcvcs.cgi/?root=Apache-SVN&view=rev&rev=1299807> Files : *
- /hadoop/common/trunk/hadoop-hdfs-project/hadoop-hdfs/CHANGES.txt *
- /hadoop/common/trunk/hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/common/HdfsServerConstants.java *
- /hadoop/common/trunk/hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/FsImage.java *
- /hadoop/common/trunk/hadoop-hdfs-project/hadoop-hdfs

```

hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/FSNamesystem.java *
/hadoop/common/trunk/hadoop-hdfs-project/hadoop-
hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/NNStorage.java *
/hadoop/common/trunk/hadoop-hdfs-project/hadoop-
hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/NameNode.java *
/hadoop/common/trunk/hadoop-hdfs-project/hadoop-
hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/NameNodeRpcServer.java *
/hadoop/common/trunk/hadoop-hdfs-project/hadoop-
hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/SecondaryNameNode.java *
/hadoop/common/trunk/hadoop-hdfs-project/hadoop-
hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/TransferFsImage.java *
/hadoop/common/trunk/hadoop-hdfs-project/hadoop-
hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/ha/BootstrapStandby.java *
/hadoop/common/trunk/hadoop-hdfs-project/hadoop-
hdfs/src/main/java/org/apache/hadoop/hdfs/tools/DFSck.java * /hadoop/common/trunk/hadoop-hdfs-
project/hadoop-hdfs/src/test/java/org/apache/hadoop/hdfs/server/namenode/TestEditLog.java *
/hadoop/common/trunk/hadoop-hdfs-project/hadoop-
hdfs/src/test/java/org/apache/hadoop/hdfs/server/namenode/ha/TestBootstrapStandby.java *
/hadoop/common/trunk/hadoop-hdfs-project/hadoop-
hdfs/src/test/java/org/apache/hadoop/test/GenericTestUtils.java

```

33. Integrated in Hadoop-Hdfs-0.23-Build #196 (See [<https://builds.apache.org/job/Hadoop-Hdfs-0.23-Build/196/>]) HDFS-2731. Add command to bootstrap the Standby Node's name directories from the Active NameNode. Contributed by Todd Lipcon. (Revision 1299809) Result = FAILURE todd : <http://svn.apache.org/viewcvcs.cgi/?root=Apache-SVN&view=rev&rev=1299809> Files : *
- ```

/hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-hdfs/CHANGES.txt *
/hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-
hdfs/src/main/java/org/apache/hadoop/hdfs/server/common/HdfsServerConstants.java *
/hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-
hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/FsImage.java *
/hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-
hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/FSNamesystem.java *
/hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-
hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/NNStorage.java *
/hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-
hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/NameNode.java *
/hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-
hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/NameNodeRpcServer.java *
/hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-
hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/SecondaryNameNode.java *
/hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-
hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/TransferFsImage.java *
/hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-
hdfs/src/main/java/org/apache/hadoop/hdfs/server/namenode/ha/BootstrapStandby.java *
/hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-
hdfs/src/main/java/org/apache/hadoop/hdfs/tools/DFSck.java * /hadoop/common/branches/branch-
0.23/hadoop-hdfs-project/hadoop-
hdfs/src/test/java/org/apache/hadoop/hdfs/server/namenode/TestEditLog.java *
/hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-
hdfs/src/test/java/org/apache/hadoop/hdfs/server/namenode/ha/TestBootstrapStandby.java *
/hadoop/common/branches/branch-0.23/hadoop-hdfs-project/hadoop-
hdfs/src/test/java/org/apache/hadoop/test/GenericTestUtils.java

```