

Item 254

git_comments:

1. nothing to do
2. ^((TOK_ALTERTABLE_ADDPARTS identifier ifNotExists?
alterStatementSuffixAddPartitionsElement+)

git_commits:

1. **summary:** HIVE-5122: Add partition for multiple partition ignores locations for non-first partitions (Navis via Thejas Nair)
message: HIVE-5122: Add partition for multiple partition ignores locations for non-first partitions (Navis via Thejas Nair) git-svn-id: <https://svn.apache.org/repos/asf/hive/branches/branch-0.12@1524940>
13f79535-47bb-0310-9956-ffa450edef68

github_issues:

github_issues_comments:

github_pulls:

github_pulls_comments:

github_pulls_reviews:

jira_issues:

1. **summary:** Add partition for multiple partition ignores locations for non-first partitions
description: <http://www.mail-archive.com/user@hive.apache.org/msg09151.html> When multiple partitions are being added in single alter table statement, the location for first partition is being used as the location of all partitions.

jira_issues_comments:

1. navis requested code review of "HIVE-5122 [jira] Add partition for multiple partition ignores locations for non-first partitions". Reviewers: JIRA HIVE-5122 Add partition for multiple partition ignores locations for non-first partitions <http://www.mail-archive.com/user@hive.apache.org/msg09151.html> TEST PLAN EMPTY REVISION DETAIL <https://reviews.facebook.net/D12411> AFFECTED FILES
ql/src/java/org/apache/hadoop/hive/ql/parse/DDLSemanticAnalyzer.java
ql/src/java/org/apache/hadoop/hive/ql/parse/HiveParser.g
ql/src/java/org/apache/hadoop/hive/ql/plan/AddPartitionDesc.java
ql/src/java/org/apache/hadoop/hive/ql/plan/DDLWork.java
ql/src/test/queries/clientpositive/add_part_exist.q ql/src/test/results/clientpositive/add_part_exist.q.out
MANAGE HERALD RULES <https://reviews.facebook.net/herald/view/differential/> WHY DID I GET THIS EMAIL? <https://reviews.facebook.net/herald/transcript/29649/> To: JIRA, navis
2. {color:red}Overall{color}: -1 at least one tests failed Here are the results of testing the latest attachment: <https://issues.apache.org/jira/secure/attachment/12598904/HIVE-5122.D12411.1.patch>
{color:red}ERROR:{color} -1 due to 1 failed/errored test(s), 2885 tests executed *Failed tests:*
{noformat} org.apache.hadoop.hive.cli.TestCliDriver.testCliDriver_create_view_partitioned {noformat}
Test results: <https://builds.apache.org/job/PreCommit-HIVE-Build/487/testReport> Console output:
<https://builds.apache.org/job/PreCommit-HIVE-Build/487/console> Messages: {noformat} Executing
org.apache.hive.ptest.execution.PrePhase Executing org.apache.hive.ptest.execution.ExecutionPhase
Executing org.apache.hive.ptest.execution.ReportingPhase Tests failed with: TestsFailedException: 1 tests
failed {noformat} This message is automatically generated.
3. navis updated the revision "HIVE-5122 [jira] Add partition for multiple partition ignores locations for non-first partitions". Fixed test fail Reviewers: JIRA REVISION DETAIL
<https://reviews.facebook.net/D12411> CHANGE SINCE LAST DIFF
<https://reviews.facebook.net/D12411?vs=38433&id=38499#toc> AFFECTED FILES
ql/src/java/org/apache/hadoop/hive/ql/parse/DDLSemanticAnalyzer.java
ql/src/java/org/apache/hadoop/hive/ql/parse/HiveParser.g

- ql/src/java/org/apache/hadoop/hive/ql/plan/AddPartitionDesc.java
ql/src/java/org/apache/hadoop/hive/ql/plan/DDLWork.java
ql/src/test/queries/clientpositive/add_part_exist.q ql/src/test/results/clientpositive/add_part_exist.q.out
ql/src/test/results/clientpositive/create_view_partitioned.q.out To: JIRA, navis
4. **body:** [~navis] Your changes in this patch make the logic in both antlr grammar and DDLSemanticAnalyzer code much cleaner, it looks good. But I didn't understand what was causing this bug (and how the change fixes it). Was it some antlr weird behavior with old grammar that caused the issue?
- label:** code-design
5. Just for the record, the AST for the plan without this fix clearly shows the problem (the location 'A' of first partition is associated with different partitions) - {code} hive> explain ALTER TABLE add_part_test ADD IF NOT EXISTS > PARTITION (ds='2010-01-01') location 'A' > PARTITION (ds='2010-02-01') location 'B' > PARTITION (ds='2010-03-01') > PARTITION (ds='2010-04-01') location 'C'; OK
ABSTRACT SYNTAX TREE: (TOK_ALTERTABLE_ADDPARTS add_part_test TOK_IFNOTEXISTS (TOK_PARTSPEC (TOK_PARTVAL ds '2010-01-01')) (TOK_PARTITIONLOCATION 'A') (TOK_PARTSPEC (TOK_PARTVAL ds '2010-02-01')) (TOK_PARTITIONLOCATION 'A') (TOK_PARTSPEC (TOK_PARTVAL ds '2010-03-01')) (TOK_PARTITIONLOCATION 'A') (TOK_PARTSPEC (TOK_PARTVAL ds '2010-04-01')) (TOK_PARTITIONLOCATION 'A')) STAGE DEPENDENCIES: Stage-0 is a root stage Stage-1 is a root stage Stage-2 is a root stage Stage-3 is a root stage STAGE PLANS: Stage: Stage-0 Stage: Stage-1 Stage: Stage-2 Stage: Stage-3 {code}
6. This looks like something we should get in for hive 0.12.
7. +1 Thanks for inputs from Harish on why the old grammar would have caused problems.
8. Thanks for the good explanation [~thejas]. Should the annotation in 'AddPartitionDesc' (the 'Path') be removed or replaced with 'Location' not be shown in test result?
9. **body:** Yes, I agree, using Location and letting it get masked would be good in general. But if the partition location is masked the existing test case will not test if the issue of this jira has been fixed, the test case would need some changes.
- label:** test
10. navis updated the revision "HIVE-5122 [jira] Add partition for multiple partition ignores locations for non-first partitions". Rebased to trunk & addressed comment (Path to Location, which is filtered by QTestUtil) Reviewers: JIRA REVISION DETAIL <https://reviews.facebook.net/D12411> CHANGE SINCE LAST DIFF <https://reviews.facebook.net/D12411?vs=38499&id=40029#toc> AFFECTED FILES
ql/src/java/org/apache/hadoop/hive/ql/parse/DDLSemanticAnalyzer.java
ql/src/java/org/apache/hadoop/hive/ql/parse/HiveParser.g
ql/src/java/org/apache/hadoop/hive/ql/plan/AddPartitionDesc.java
ql/src/java/org/apache/hadoop/hive/ql/plan/DDLWork.java
ql/src/test/queries/clientpositive/add_part_exist.q ql/src/test/results/clientpositive/add_part_exist.q.out
ql/src/test/results/clientpositive/create_view_partitioned.q.out To: JIRA, navis
11. Navis, as I mention in my previous comment, with masking the test no longer checks if the locations are being picked up correctly. Ie, we won't know if someone introduces a bug that causes same problem again, and causes the first location to be associated with each of the partitions. The test case needs to be changed. Maybe, add partitions with data and use a select query selecting data from the partitions, to verify that correct locations are associated with the partitions?
12. Good idea! I'll update the patch.
13. navis updated the revision "HIVE-5122 [jira] Add partition for multiple partition ignores locations for non-first partitions". Addressed comments Reviewers: JIRA REVISION DETAIL <https://reviews.facebook.net/D12411> CHANGE SINCE LAST DIFF <https://reviews.facebook.net/D12411?vs=40029&id=40059#toc> AFFECTED FILES
ql/src/java/org/apache/hadoop/hive/ql/parse/DDLSemanticAnalyzer.java
ql/src/java/org/apache/hadoop/hive/ql/parse/HiveParser.g
ql/src/java/org/apache/hadoop/hive/ql/plan/AddPartitionDesc.java
ql/src/java/org/apache/hadoop/hive/ql/plan/DDLWork.java
ql/src/test/org/apache/hadoop/hive/ql/QTestUtil.java ql/src/test/queries/clientpositive/add_part_multiple.q ql/src/test/results/clientpositive/add_part_multiple.q.out
ql/src/test/results/clientpositive/create_view_partitioned.q.out To: JIRA, navis
14. Looks good. +1
15. Resubmitting patch for kicking off precommit tests on updated patch.
16. Copy of HIVE-5122.D12411.2.patch to get pre-commit tests running

17. I previously copied wrong file. This time uploading D12411.4.patch as HIVE-5122.4.patch (conforms to format expected by <https://cwiki.apache.org/confluence/display/Hive/Hive+PreCommit+Patch+Testing>)
18. {color:green}Overall{color}: +1 all checks pass Here are the results of testing the latest attachment: <https://issues.apache.org/jira/secure/attachment/12604159/HIVE-5122.4.patch> {color:green}SUCCESS: {color} +1 3128 tests passed Test results: <https://builds.apache.org/job/PreCommit-HIVE-Build/832/testReport> Console output: <https://builds.apache.org/job/PreCommit-HIVE-Build/832/console> Messages: {noformat} Executing org.apache.hive.ptest.execution.PrePhase Executing org.apache.hive.ptest.execution.ExecutionPhase Executing org.apache.hive.ptest.execution.ReportingPhase {noformat} This message is automatically generated.
19. Patch committed to trunk and 0.12 branch. Thanks for the contribution Navis!
20. FAILURE: Integrated in Hive-trunk-hadoop2-ptest #107 (See [<https://builds.apache.org/job/Hive-trunk-hadoop2-ptest/107/>]) HIVE-5122: Add partition for multiple partition ignores locations for non-first partitions (Navis via Thejas Nair) (thejas: <http://svn.apache.org/viewcvcs.cgi/?root=Apache-SVN&view=rev&rev=1524930>) *
- /hive/trunk/ql/src/java/org/apache/hadoop/hive/ql/parse/DDLSemanticAnalyzer.java *
 - /hive/trunk/ql/src/java/org/apache/hadoop/hive/ql/parse/HiveParser.g *
 - /hive/trunk/ql/src/java/org/apache/hadoop/hive/ql/plan/AddPartitionDesc.java *
 - /hive/trunk/ql/src/java/org/apache/hadoop/hive/ql/plan/DDLWork.java *
 - /hive/trunk/ql/src/test/org/apache/hadoop/hive/ql/QTestUtil.java *
 - /hive/trunk/ql/src/test/queries/clientpositive/add_part_multiple.q *
 - /hive/trunk/ql/src/test/results/clientpositive/add_part_multiple.q.out *
 - /hive/trunk/ql/src/test/results/clientpositive/create_view_partitioned.q.out
21. FAILURE: Integrated in Hive-trunk-hadoop1-ptest #174 (See [<https://builds.apache.org/job/Hive-trunk-hadoop1-ptest/174/>]) HIVE-5122: Add partition for multiple partition ignores locations for non-first partitions (Navis via Thejas Nair) (thejas: <http://svn.apache.org/viewcvcs.cgi/?root=Apache-SVN&view=rev&rev=1524930>) *
- /hive/trunk/ql/src/java/org/apache/hadoop/hive/ql/parse/DDLSemanticAnalyzer.java *
 - /hive/trunk/ql/src/java/org/apache/hadoop/hive/ql/parse/HiveParser.g *
 - /hive/trunk/ql/src/java/org/apache/hadoop/hive/ql/plan/AddPartitionDesc.java *
 - /hive/trunk/ql/src/java/org/apache/hadoop/hive/ql/plan/DDLWork.java *
 - /hive/trunk/ql/src/test/org/apache/hadoop/hive/ql/QTestUtil.java *
 - /hive/trunk/ql/src/test/queries/clientpositive/add_part_multiple.q *
 - /hive/trunk/ql/src/test/results/clientpositive/add_part_multiple.q.out *
 - /hive/trunk/ql/src/test/results/clientpositive/create_view_partitioned.q.out
22. FAILURE: Integrated in Hive-trunk-h0.21 #2345 (See [<https://builds.apache.org/job/Hive-trunk-h0.21/2345/>]) HIVE-5122: Add partition for multiple partition ignores locations for non-first partitions (Navis via Thejas Nair) (thejas: <http://svn.apache.org/viewcvcs.cgi/?root=Apache-SVN&view=rev&rev=1524930>) *
- /hive/trunk/ql/src/java/org/apache/hadoop/hive/ql/parse/DDLSemanticAnalyzer.java *
 - /hive/trunk/ql/src/java/org/apache/hadoop/hive/ql/parse/HiveParser.g *
 - /hive/trunk/ql/src/java/org/apache/hadoop/hive/ql/plan/AddPartitionDesc.java *
 - /hive/trunk/ql/src/java/org/apache/hadoop/hive/ql/plan/DDLWork.java *
 - /hive/trunk/ql/src/test/org/apache/hadoop/hive/ql/QTestUtil.java *
 - /hive/trunk/ql/src/test/queries/clientpositive/add_part_multiple.q *
 - /hive/trunk/ql/src/test/results/clientpositive/add_part_multiple.q.out *
 - /hive/trunk/ql/src/test/results/clientpositive/create_view_partitioned.q.out
23. FAILURE: Integrated in Hive-trunk-hadoop2 #444 (See [<https://builds.apache.org/job/Hive-trunk-hadoop2/444/>]) HIVE-5122: Add partition for multiple partition ignores locations for non-first partitions (Navis via Thejas Nair) (thejas: <http://svn.apache.org/viewcvcs.cgi/?root=Apache-SVN&view=rev&rev=1524930>) *
- /hive/trunk/ql/src/java/org/apache/hadoop/hive/ql/parse/DDLSemanticAnalyzer.java *
 - /hive/trunk/ql/src/java/org/apache/hadoop/hive/ql/parse/HiveParser.g *
 - /hive/trunk/ql/src/java/org/apache/hadoop/hive/ql/plan/AddPartitionDesc.java *
 - /hive/trunk/ql/src/java/org/apache/hadoop/hive/ql/plan/DDLWork.java *
 - /hive/trunk/ql/src/test/org/apache/hadoop/hive/ql/QTestUtil.java *
 - /hive/trunk/ql/src/test/queries/clientpositive/add_part_multiple.q *
 - /hive/trunk/ql/src/test/results/clientpositive/add_part_multiple.q.out *
 - /hive/trunk/ql/src/test/results/clientpositive/create_view_partitioned.q.out

24. This issue has been fixed and released as part of 0.12 release. If you find further issues, please create a new jira and link it to this one.