

**git\_comments:**

1. Get max possible key value

**git\_commits:**

1. **summary:** PHOENIX-992: Replace `hTable.getRowOrBefore` with `ReverseScan` to get `maxKey` in `StatsManagerImpl`  
**message:** PHOENIX-992: Replace `hTable.getRowOrBefore` with `ReverseScan` to get `maxKey` in `StatsManagerImpl`

**github\_issues:**

**github\_issues\_comments:**

**github\_pulls:**

**github\_pulls\_comments:**

**github\_pulls\_reviews:**

**jira\_issues:**

1. **summary:** Replace `hTable.getRowOrBefore` with `ReverseScan` to get `maxKey` in `StatsManagerImpl`  
**description:** Below is the java doc of the function in HBase and we should remove it ASAP {noformat}  
\* @deprecated As of version 0.92 this method is deprecated without \* replacement. \* `getRowOrBefore` is used internally to find entries in `hbase:meta` and makes \* various assumptions about the table (which are true for `hbase:meta` but not \* in general) to be efficient. {noformat}

**jira\_issues\_comments:**

1. +1 on this idea, as our dependency on this method is minimal as you've pointed out. Would this be for 4.0 only, or is this broken for 0.94 too? Should we file an HBase bug as well?
2. [~giacomotaylor] The issue may in 3.0 as well but my understanding that `ReverseScan` isn't in 0.94 branch so there is no good way to replace it. Here is the patch which is using reverse scan to get the `maxKey` of a table.
3. +1
4. **body:** `getRowOrBefore` shouldn't be used and has been deprecated long enough, HBase should yank it from the API, but I wouldn't expect that to happen in 0.94.  
**label:** code-design
5. Thanks [~giacomotaylor] for the review! All tests passed locally with small tweaks on the number of parallel scans(more parallels) which may mean previous `maxKey` is smaller. If no objections, I'll commit v2 patch tomorrow. Thanks.
6. +1
7. There is one difference though wrt performance (in using reverse scan vs `HTable.getRowOrBefore()`) The `HTable.getRowOrBefore()` will do one parse of KVs within a HFile block. The use of reverse scan(as in the patch) will do parse within an HFile block 2 times. 1st time for the seek to last row when the reverse scanner is opened. The next() call will return the last KV but will a move to the previous KV. This in effect will do a parse of KVs in the HFile block again. Just telling the difference.
8. +1. Thanks!
9. Thanks for all the review & comments! I've integrated the fix into branch 4.0 & master.
10. FAILURE: Integrated in Phoenix-hadoop2 #2 (See [<https://builds.apache.org/job/Phoenix-hadoop2/2/>]) PHOENIX-992: Replace `hTable.getRowOrBefore` with `ReverseScan` to get `maxKey` in `StatsManagerImpl` (jzhong: rev dff13fe6c16da9e6f70421b292bf64a0422b4ebf) \* phoenix-core/src/main/java/org/apache/phoenix/query/StatsManagerImpl.java \* phoenix-core/src/it/java/org/apache/phoenix/end2end/TenantSpecificViewIndexIT.java \* phoenix-core/src/it/java/org/apache/phoenix/end2end/BaseTenantSpecificViewIndexIT.java
11. SUCCESS: Integrated in Phoenix #228 (See [<https://builds.apache.org/job/Phoenix/228/>]) PHOENIX-992: Replace `hTable.getRowOrBefore` with `ReverseScan` to get `maxKey` in `StatsManagerImpl` (jzhong:

rev dff13fe6c16da9e6f70421b292bf64a0422b4ebf) \* phoenix-  
core/src/it/java/org/apache/phoenix/end2end/BaseTenantSpecificViewIndexIT.java \* phoenix-  
core/src/main/java/org/apache/phoenix/query/StatsManagerImpl.java \* phoenix-  
core/src/it/java/org/apache/phoenix/end2end/TenantSpecificViewIndexIT.java  
12. Bulk close of all issues that has been resolved in a released version.