

Item 367

git_comments:

git_commits:

1. **summary:** ARROW-252: Add implementation guidelines to the documentation
message: ARROW-252: Add implementation guidelines to the documentation Author: Julien Le Dem <julien@dremio.com> Closes #120 from julienledem/arrow_252_impl_guidelines and squashes the following commits: caf6994 [Julien Le Dem] ARROW-252: review feedback 6b68ce1 [Julien Le Dem] ARROW-252: Add implementation guidelines to the documentation

github_issues:

github_issues_comments:

github_pulls:

1. **title:** ARROW-252: Add implementation guidelines to the documentation
body:

github_pulls_comments:

1. Looks good to me. Would be great for someone else to take a look too

github_pulls_reviews:

1. Mention whether custom metadata are in/out of scope for V1 (or ever). I would initially vote for out of scope to keep things simple.
2. minor. Capitalize "any", add period at the end.
3. done
4. done

jira_issues:

1. **summary:** Add implementation guidelines to the documentation
description: I'd like to add a paragraph to the documentation providing implementation guidelines: An execution engine (or framework, or UDF executor, or storage engine, etc) can use only a subset of the arrow Arrow spec or extend it given the following constraints: Implementing a subset: 1) If it is only producing (and not consuming) arrow vectors. - any subset of the vector spec and the corresponding metadata can be implemented 2) If it is consuming *and* producing vectors - there is a minimal subset of vectors to be supported (To Be Defined) - production of a subset of vectors and their corresponding metadata is always fine - consumption of vectors should at least convert the unsupported input vectors to the supported subset (for example Timestamp.millis to timestamp.micros or int32 to int64) An execution engine implementor can also extend their memory representation with their own vectors internally as long as they are never exposed. Before sending data to another system expecting Arrow data these custom vectors should be converted to a type that exist in the Arrow spec. An exemple of this is operating on compressed data.

jira_issues_comments:

1. +1 I agree with these sentiments
2. +1 as well
3. PR: <https://github.com/apache/arrow/pull/120>
4. Issue resolved by pull request 120 [<https://github.com/apache/arrow/pull/120>]